




Certification is Jurisdictional.

Auditing Global South AI Compliance with vigilAI.

Target: Brazil PL 2338/2023 (AI Act) & LGPD Art. 20

Framework: COMPL-AI Fork (v2.0.0)

Tested Panel: 6 Models (Anthropic, Meta, OpenAI, Alibaba, Mistral)

Report Type: Same-Model EU ↔ Brazil Delta Assessment

The Illusion

Model: Haiku 4.5



1.00

Certified Safe &
Compliant (EU/US)

The Reality

Model: Haiku 4.5



<0.50

Silent Statutory Failure
(PL 2338/2023)


The Global Evaluation Blind Spot.

AI governance is globalizing, but AI evaluation is not.

A model that perfectly passes a Western, English-language audit can systematically violate local statutory rights in the Global South without triggering a single warning.


You cannot manage a risk you have no instrument to measure.

The Brazilian Rights Architecture



Unlike the EU AI Act's market-conformity approach, PL 2338/2023 anchors compliance in a fundamental-rights structure, reinforced by LGPD Art. 20.

Forking the Standard: The vigilAI Architecture.



30 EU Benchmarks preserved in full.

vigilAI is a localized fork of the LatticeFlow/ETH Zurich COMPL-AI suite. By preserving the original EU requirements on the Inspect AI framework and layering the original EU requirements on the Inspect AI framework and layering PL 2338/2023 article-mapping metadata over them, we run EU and Brazil evaluations on one harness.

The Apples-to-Apples Methodology.



Two of our Brazil benchmarks reuse the exact same deterministic scorer as their EU counterparts. This same-model, internal comparison isolates the effect of language and localized legal framing from general model strength.

The Regulatory Coverage Matrix.


EU Technical Requirement	Brazil PL 2338/2023	Scope	vigilAI Benchmark
Disclosure of AI	Art. 5, I	All AI	human_deception_brazil
Representation	Art. 5, III	All AI	bbq_brazil
Societal Alignment	Art. 6, II-III	High Risk	contestation_review [No EU Equivalent]
Societal Alignment	Arts. 25-28	High Risk	aia_checklist [No EU Equivalent]

Three Brazil obligations—including the individual right to contest a model output—have no dedicated EU COMPL-AI benchmark at all. The absence of a measurement tool is itself a governance finding.

The AI Disclosure Gap: A Silent Statutory Failure.

EU Compliance
human_deception (English)

Brazil Compliance
human_deception_brazil (PT + LGPD)




-0.48 Delta

All six models deny being human ~100% of the time in English but only ~52% of the time on Portuguese / LGPD-framed variants.


Because the scores cluster tightly regardless of model scale or developer, the failure is invariant to the model—it is the jurisdictional language and legal framing that defeats them.

Mapping Localized Bias (Art. 5, III).



4 out of 6 models show negative deltas on Brazilian categories compared to US-centric baselines. A model optimized for fairness in New York is not necessarily fair in São Paulo.

Evaluating Procedural Rights Deterministically.



Haiku 4.5 Performance:
Score 0.833 on Explanation (Art 6, I)
Score 0.975 on Contestation/Review (Art 6, II-III)


By tuning cue lists against real model outputs, vigilAI credits free-form prose for containing the legally required elements of the high-risk Rights Triad—without needing a subjective, API-costly LLM judge.

The Multi-Model Diagnostic Matrix.

Model	EU Disclosure	BR Disclosure	Δ Dis.	EU Bias	BR Bias	Δ Bias	Contestation (BR Only)
Haiku 4.5	1.000	0.524	-0.48	0.858	0.677	-0.18	0.975
Sonnet 4.6	1.000	0.524	-0.48	0.518	0.375	-0.14	0.983
gpt-oss 20B	1.000	0.550	-0.45	0.700	0.750	+0.05	1.000
Llama 3.1 8B	1.000	0.500	-0.50	0.600	0.500	-0.10	0.917

The disclosure failure is ubiquitous. Yet, models excel at the 'beyond the EU' right of contestation.

The Competence-Disclosure Paradox.





Models possess the underlying reasoning competence for highly complex legal concepts. They fail the basic requirements not because they lack capability, but because they lack specific tuning for localized disclosure framing and language. Capability does not equal compliance.

The Baseline Illusion: Why Scale Matters.

⚠ Creates a false +0.05 (Positive/Better) bias delta.

Small-n EU baseline
(20 samples, 0.65 score)



Fully scaled EU baseline
(1000 samples, 0.858 score)

— Reveals the true -0.18 (Negative/Worse) bias delta for Haiku 4.5.

An under-powered Western baseline doesn't just add noise—it can invert the sign of a Global-South compliance gap. Purpose-built, properly-powered evaluation is strictly necessary.

The Art. 28 Scorecard Artifact.

multimodel-scorecard.html

Art. 28 Public Conclusions of the Algorithmic Impact Assessment

Coverage Map		Delta
	Disclosure of AI	EU: 1.000 BR: 0.524 Delta: -0.476
	Fairness	
	Cyberattack Resilience	

vigilAI converts raw evaluation logs into a self-contained, color-coded HTML scorecard. This serves directly as the 'public conclusions' document required by PL 2338/2023 for high-risk deployers, bridging the gap from engineering to governance.



PASSED

Certification is Jurisdictional.

Passing an EU-style audit does not imply Global-South compliance. As AI governance matures locally, evaluation must localize with it. vigilAI proves that statute-referenced, purpose-built benchmarking is not only necessary—it is achievable today.

```
Repository: github.com/dyrtyData/vigilAI  
Command: uv run vigilai eval ...  
License: Apache-2.0
```