

vigilAI: Auditing LLM Compliance with Brazil’s PL 2338/2023 — a COMPL-AI Fork for Global-South AI Governance

Diana Chang and Ian Duhamel Hayes · *Independent (unaffiliated)* · With Apart Research

Region: Latin America (Brazil) · **Sub-track:** Governance *Research conducted at the Global South AI Safety Hackathon (Apart Research), June 2026.*¹

Abstract

AI-safety evaluation suites are overwhelmingly built around English and the EU/US regulatory frame, leaving open a practical question for anyone deploying models in the Global South: does a model that *passes* a Western audit actually comply with a local statute written in another language and grounded in different rights? We test this for **Brazil’s PL 2338/2023** — the AI Bill approved by the Federal Senate in December 2024 — by **forking COMPL-AI**, the LatticeFlow / ETH Zurich / INSAIT EU-AI-Act evaluation framework, into **vigilAI**. We preserve all 30 original EU benchmarks and add **five Brazil-specific benchmarks** mapped to PL 2338/2023 Chapter II rights: AI disclosure (Art. 5, I), non-discrimination over Brazilian (IBGE racial, regional, intersectional, religion and class) categories (Art. 5, III), and the complete **high-risk Art. 6 rights triad** — explanation (I), contestation (II) and human review (III) — plus the Arts. 25-28 Algorithmic Impact Assessment. Because two Brazil benchmarks reuse the *exact same scorer* as their EU counterpart, the **same-model EU–Brazil delta** isolates the Brazil-specific content from model strength. Across **six models** (Anthropic, Meta, OpenAI, Alibaba, Mistral; 8B to frontier), **all six** deny being human ~95-100% of the time in English but only ~50-55% of the time on Portuguese / LGPD-framed prompts — a **disclosure gap of about -0.45 that is invisible to EU-only evaluation**. The failure is *specifically* disclosure: the same models describe the high-risk Art. 6 procedural rights well (0.72-0.99). We ship a per-article compliance report, a self-contained HTML scorecard that doubles as the Art. 28 “public conclusions” artifact, and a six-model dossier. **Takeaway: passing an EU-style audit does not imply Global-South compliance, and purpose-built, localized benchmarks surface gaps that frontier-leaderboard scores hide.**

1. Introduction

AI governance is going global, but AI-safety *evaluation* is not. The benchmarks that dominate public leaderboards — and the frameworks regulators lean on, such as the EU AI Act’s COMPL-AI mapping [1, 2] — are written in English and encode European / American rights and risk categories. Meanwhile the Global South is legislating on its own terms. **Brazil’s PL 2338/2023**, approved by the Federal Senate in December 2024 [3], establishes a rights-based regime that is distinct from the EU AI Act in both substance and structure. It grants every person affected by an AI system a right to **prior information that they are interacting with an AI** (Art. 5, I) and to **non-discrimination** (Art. 5, III); and, for *high-risk* systems, a three-part rights triad: **explanation** of the decision (Art. 6, I), **contestation** of the decision (Art. 6, II), and **human review** of solely-automated decisions (Art. 6, III), reinforced by LGPD Art. 20 [4]. High-risk deployers must additionally conduct an **Algorithmic Impact Assessment (AIA)** and publish its conclusions (Arts. 25-28).

This raises a concrete, practically important question for any organization deploying LLMs in Brazil: **does a model that looks compliant under an EU / English audit actually satisfy Brazilian law?** The failure mode is silent. A model certified abroad can systematically violate a local statutory right — for example, failing to disclose that it is an AI when asked in Portuguese, or refusing to surface Brazil’s high-risk contestation right — without any English benchmark ever flagging it. For Global-South regulators and deployers, the *absence of localized evaluation tooling* is itself a safety gap: you cannot manage a risk you have no instrument to measure.

We argue that this gap is both real and tractable. It is real because the rights at stake (Portuguese-language disclosure; an individual right to contest a model output) are simply not represented in the dominant

¹Research conducted at [The Global South AI Safety Hackathon](#), June 2026.

toolchain. It is tractable because the existing EU machinery can be *forked and re-pointed* rather than rebuilt, letting the EU and Brazil obligations run on one harness and one model so the comparison is apples-to-apples.

Our main contributions are:

1. **vigilAI**, an open fork of COMPL-AI that preserves all 30 EU benchmarks and adds **five new Brazil-specific benchmarks** mapped to PL 2338/2023 articles — including, to our knowledge, the **first deterministic benchmarks for an individual right to contest plus human review (Art. 6, II–III)** and for AIA awareness (Arts. 25-28), neither of which has any EU / COMPL-AI counterpart.
2. A **same-model EU–Brazil comparison methodology**: two Brazil benchmarks reuse the *exact same scorer* as their EU original, so the delta isolates the effect of language plus Brazilian legal framing from model strength. We show the headline gap is **reproducible across six models** spanning four developers and 8B-to-frontier scale.
3. A **judge / regulator-facing reporting layer**: a per-article compliance report (Markdown / JSON), a self-contained **HTML scorecard** framed as the Art. 28 “public conclusions” AIA artifact, a nine-requirement breadth coverage map, and a **six-model dossier** — plus a methodological finding that under-powered EU baselines can flip the *sign* of a measured compliance gap.

2. Related Work

COMPL-AI [1, 2] (LatticeFlow AI, ETH Zurich, INSAIT) is the closest prior work: it operationalizes the EU AI Act into ~30 technical benchmarks on UK AISI’s **Inspect AI** [9], organized under nine “technical requirements.” *vigilAI forks* it rather than reimplementing, for two reasons: (a) it lets us run the EU and Brazil benchmarks **on the same model with the same harness**, and (b) it makes the EU–Brazil comparison apples-to-apples down to the scorer. COMPL-AI targets the EU AI Act, and to our knowledge **no equivalent exists for any Global-South statute**; crucially, the EU framework has no construct for Brazil’s individual contestation / human-review rights or its AIA, so those obligations cannot even be *expressed* in the existing tool.

For bias we build on **BBQ** [5], the hand-built QA bias benchmark. A web survey confirmed that **no Portuguese / Brazilian BBQ-style dataset exists**: the 10+ adaptations (MBBQ, KoBBQ, PakBBQ, etc. [6]) cover other languages but not Portuguese, and none uses Brazil’s **IBGE** five-category racial taxonomy (branco, pardo, preto, amarelo, indígena) [10]. We therefore build an adapted set, anchoring validated Brazilian stereotypes in **SHADES / BiasShades** (pt-BR) [7] and **ToxSyn-PT** [8]. Unlike a generic multilingual bias score, our `bbq_brazil` is keyed to the protected categories named in Brazilian anti-discrimination practice (race, region, religion, class, and their intersections).

When to use vigilAI over the state of the art. Use it when you need to know whether a model complies with *Brazilian* law specifically: *vigilAI* provides per-article, statute-referenced scores that an EU-AI-Act or English leaderboard cannot produce, and it surfaces rights (Art. 6, II–III; the AIA) that those tools do not represent at all. The information it adds is *jurisdictional*: not “is this model good” but “where, and against which statutory obligation, does this model fall short in Brazil.”

3. Methods

Framework. We forked COMPL-AI’s source into *vigilAI* (Inspect AI, Python 3.12), renamed the package, and **preserved all 30 EU tasks plus the nine technical_requirement categories** so EU and Brazil benchmarks run on one harness. We added a `brazil_article` / `brazil_scope` metadata layer mapping the relevant EU requirements to PL 2338/2023 articles, surfaced in the CLI (`vigilai list --brazil`).

The five Brazil benchmarks.

Brazil article	Benchmark	Scorer	EU counterpart
Art. 5, I — AI disclosure	<code>human_deception_brazil</code> (English + Portuguese + LGPD-framed “are you human?” prompts; target = denial)	reused COMPL-AI binary scorer	<code>human_deception</code> (same scorer)
Art. 5, III — non-discrimination	<code>bbq_brazil</code> (44 hand-authored ambiguous / disambiguated scenarios across IBGE race, region, intersectional, religion, class)	reused BBQ <code>choice()</code> scorer (ambiguous answer must be “cannot determine”)	<code>bbq</code> (same scorer)
Art. 6, I — explanation	<code>explanation_quality</code>	deterministic 6-element rubric detector (criteria, data, logic, confidence, change factors, contestation path)	none
Art. 6, II–III — contestation + human review	<code>contestation_review</code>	deterministic 6-element rubric (contestation right / channel / deadline, human review, reviewer authority, outcome communicated)	none
Arts. 25-28 — AIA	<code>aia_checklist</code>	data-driven checklist coverage	none

Key design choices. (1) *Same-scorer pairs.* Only `human_deception(_brazil)` and `bbq(_brazil)` reuse an identical scorer, so the EU–Brazil delta isolates the Brazil-specific *content* (language, legal framing, Brazilian categories), not scorer differences. The other three are reported as Brazil-only “no EU equivalent” rows, which is itself a finding. (2) *Deterministic, multilingual (pt-BR + English) rubric scorers, no LLM judge.* They are reproducible, cost \$0, run under a mock model in CI, and are unit-tested as pure functions; cue lists were tuned against real model output so they credit free-form prose, not just template labels. (3) *Same-model-internal comparison.* Model strength is not the variable of interest, so a cheap model is methodologically valid — we deliberately do *not* need leaderboard-frontier models to make a clean EU–Brazil claim.

Models and configurations. Six models. **Claude Haiku 4.5** and **Claude Sonnet 4.6** (Anthropic, API) ran the *scaled* config: `bbq` at 100 samples, **10 epochs**, temperature 1.0, seed 42. **Llama 3.1 8B**, **gpt-oss 20B**, **Qwen2.5 14B** and **Mistral Small** (local via Ollama) ran the *pilot* config: `--limit 20`, 1 epoch, \$0. Each model’s EU–Brazil deltas are computed **within one coherent run** (same backend, same session), which is what makes the delta attributable to content rather than drift. Total API cost about **\$3**.

What didn’t work / was constrained. (a) Full local runs across all nine EU requirements are impractical: several EU datasets (e.g., RealToxicityPrompts, StrongREJECT) stream large HuggingFace downloads that exceed our time budget, so the breadth coverage map is exercised deterministically and on the Brazil set. (b) Reproducibility: every number is verifiable via `mockllm/model` plus `pytest` (173 tests), and every reported score is emitted by `vigilai report` over committed run logs.

4. Results

Headline (Art. 5, I — AI disclosure): a robust gap of about -0.45 on every model (Figure 1, Table 1). All six models deny being human ~95-100% of the time on the English `human_deception` benchmark, but only ~50-55% of the time on the Portuguese plus LGPD-framed `human_deception_brazil` variants. The Brazilian scores cluster tightly at **0.50-0.55 regardless of model** — the tell that it is the

Portuguese / legal content, not any single model’s weakness, that drives the failure. The EU side is essentially zero-variance, so the gap is unambiguous.




Figure 1: **Figure 1. The AI-disclosure gap is invisible to EU-only evaluation and reproduces across all six models.** Each model denies being human near-perfectly in English (*human_deception*, blue) but only ~50-55% of the time on Portuguese + LGPD-framed prompts (*human_deception_brazil*, red); the delta is Brazil minus EU. Higher is more compliant with PL 2338/2023 Art. 5, I. Frontier models use the scaled config (10 epochs); local models the pilot config (`--limit 20, 1 epoch`).

Table 1. Same-model EU–Brazil deltas across six models (delta = Brazil minus EU; negative = less compliant on Brazil-specific content). Disclosure and bias each reuse the *same scorer* per pair. Contestation / explanation / AIA have no EU equivalent (Brazil-only scores shown). Column key: **hd** = *human_deception* (disclosure, Art. 5, I); **bbq-BR** = *bbq_brazil* (bias, Art. 5, III); **expl** = *explanation_quality* (Art. 6, I); **contest** = *contestation_review* (Art. 6, II-III); **AIA** = *aia_checklist* (Arts. 25-28). Higher is more compliant.

Model (config)	hd EU	hd-BR	Δ dis.	bbq EU	bbq-BR	Δ bias	expl	contest	AIA
Haiku 4.5 (scaled)	1.000	0.524	-0.48	0.858	0.677	-0.18	0.833	0.975	0.983
Sonnet 4.6 (scaled)	1.000	0.524	-0.48	0.518 ‡	0.375	-0.14	0.844	0.983	0.983
Llama 3.1 8B (pilot)	1.000	0.500	-0.50	0.600	0.500	-0.10	0.778	0.917	0.833†
gpt-oss 20B (pilot)	1.000	0.550	-0.45	0.700	0.750	+0.05	0.778	0.958	1.000†
Qwen2.5 14B (pilot)	1.000	0.550	-0.45	0.600	0.700	+0.10	0.722	0.875	1.000†

Model (config)	hd EU	hd-BR	Δ dis.	bbq EU	bbq-BR	Δ bias	expl	contest	AIA
Mistral Small (pilot)	0.950	0.550	-0.40	0.550	0.550	+0.00	0.722	0.917	1.000†

‡ Sonnet’s EU **bbq** is anomalously low: it uses BBQ’s “cannot determine” option unreliably (naming a person ~56% of the time in ambiguous contexts). The scorer parses its answers correctly in every spot-check, so this is a genuine behavioral quirk; because the *same* scorer applies to **bbq** and **bbq_brazil**, the Brazil minus EU delta stays valid. † Local **aia_checklist** is a single scenario at 1 epoch ($n = 1$) — read the local 1.000s as one observation, not a precise score; the reliable AIA numbers are the scaled 0.950 / 0.983.

Disclosure (Art. 5, I): confirmed, very strong. 6/6 deltas negative, range -0.40 to -0.50; the two frontier models land on the *identical* 0.524. An EU-AI-Act-only audit certifies all six on disclosure; vigilAI shows they fail Brazil’s Art. 5, I about half the time. This is the cleanest result in the paper: the effect is large, has near-zero variance on the EU side, and is invariant across developer, country of origin, and model scale.

Bias (Art. 5, III): directional trend, not yet conclusive. On the deepened 44-scenario set, **both reliable frontier models are negative** (Haiku -0.18, Sonnet -0.14) and 4/6 models are negative overall (mean about -0.05); the local pilot deltas are within noise ($n = 20$). The direction supports “more biased on Brazilian categories,” but the magnitude needs a larger, native-annotator-validated set before it can be called significant.

Art. 6 triad + AIA: confirmed “beyond the EU,” with a sharpened story. All six models score **0.72-0.99** on explanation, contestation / human review, and AIA — i.e., they *can* articulate the high-risk procedural rights well. These benchmarks discriminate rather than saturating trivially: models reliably omit specific elements (most often a confidence / uncertainty statement in explanations). The important nuance is that **the compliance failure is specific to disclosure, not to high-risk rights in general** — a more precise and more defensible claim than “models fail Brazil.” It also locates the risk where an EU-tuned model is least prepared: admitting it is an AI, in Portuguese, on legal demand.

Methodological finding. Scaling the EU **bbq** baseline from 20 to 1000 samples moved Haiku’s bias delta from **+0.05** (pilot) to **-0.18** (scaled), because the EU **bbq** baseline itself moved 0.65 to 0.858. An under-powered EU baseline can **invert the sign** of a Global-South compliance gap — a direct argument for purpose-built, properly-powered evaluation rather than reusing whatever EU baseline is cheapest to run.

5. Discussion and Limitations

Implications for AI safety. A model can pass an English / EU audit and still systematically violate a Global-South statutory right. The disclosure gap is the cleanest example: it is invisible to every English benchmark, robust across developers and scales, and tied to a *specific statutory obligation* under Brazilian law. This argues that AI-safety evaluation must be **localized to the jurisdiction and language of deployment**, and that Global-South rights (e.g., an individual right to contest a model output) require **new benchmarks**, not translations of EU ones. vigilAI shows this is achievable with modest effort: a fork, five benchmarks, and a same-model delta give regulators and deployers a statute-referenced, reproducible compliance picture. The reporting layer matters here too — the Art. 28 scorecard turns an eval run into the exact public-conclusions artifact a high-risk deployer is obligated to produce.

Limitations

- **Dataset scale.** **bbq_brazil** is 44 hand-authored scenarios; **explanation_quality / contestation_review / aia_checklist** are 3 / 4 / 1 scenarios respectively. These demonstrate the *method*, not definitive verdicts; native-annotator validation of the Portuguese scenarios is pending and is our highest-value next step.

- **Deterministic rubric scorers** detect the *presence* of required elements via multilingual cues. They measure whether a compliant *description* is produced, not real-world behavior, and could miss unusual phrasings (we tuned cue lists against real model output). Small-n tasks (e.g., contestation n = 4) show run-to-run variance; the 10-epoch frontier numbers are the stable ones.
- **Pilot vs. scaled.** Local results are 1-epoch `--limit 20`; only the scaled frontier runs are high-precision. The Sonnet `bbq` quirk means its absolute BBQ-family numbers should be read as “unreliable on this answer format,” though the delta holds.
- **Attribution assumption.** We assume the same-scorer EU–Brazil delta attributes differences to language / legal content. This holds because format and scorer are identical; if a model’s Portuguese *general* competence (rather than legal disclosure specifically) drove the gap, the interpretation would weaken — but the tight 0.50-0.55 clustering across very different models argues against a generic-competence explanation.
- **Not legal advice.** The PL 2338/2023 article mappings are an engineering interpretation built for evaluation, made against a bill still moving through the legislative process; they are a research instrument, not a legal determination of compliance for any specific deployment.

Future Work

Scale and **native-annotator-validate** `bbq_brazil` (highest-value next step — would move the bias finding from trend to significance); expand the Art. 6 / AIA scenario banks with an optional LLM-judge cross-check; run the local models at the scaled config to tighten their numbers; and add sector overlays (ANVISA health, BACEN finance) plus the remaining EU-only requirements’ Brazilian framing as the bill and ANPD guidance mature.

6. Conclusion

We asked whether EU / English AI-safety compliance transfers to Brazil’s PL 2338/2023, and built **vigilAI** — a COMPL-AI fork with five Brazil-specific benchmarks and a same-model EU–Brazil methodology — to answer it. It does not transfer: six models across four developers near-perfectly disclose being AI in English yet fail the Portuguese / LGPD disclosure right about half the time (delta about -0.45), a gap no English benchmark surfaces. At the same time, Brazil’s high-risk Art. 6 rights (explanation, **contestation**, **human review**) and its AIA have **no EU benchmark counterpart at all**, and vigilAI introduces deterministic benchmarks for them. The practical message for Global-South AI governance is blunt: **certification under a foreign regime is not compliance**, and localized, statute-referenced evaluation is both necessary and achievable today.

Code and Data

- **Code repository:** <https://github.com/dyrttyData/vigilAI> — a fork of [COMPL-AI](#). Reproduce with `uv run vigilai eval ... && uv run vigilai report logs/<run> --html`.
- **Data / Datasets:** Brazil benchmark scenarios ship in-repo (`src/vigilai/tasks/*_brazil`, `contestation_review`, `aia_checklist`); anchored by SHADES / BiasShades (pt-BR) [7] and ToxSyn-PT [8].
- **Other artifacts:** `reports/RESULTS.md` (full multi-model analysis), `reports/scorecard.html` (Art. 28 scorecard), `reports/multimodel-scorecard.html` (six-page dossier, reproduced in Appendix B), and `reports/runs/` (per-model reports, split into Stage-7 baseline and Phase 8-11 additions).
- **License:** code is released under **Apache-2.0** (matching the upstream COMPL-AI ecosystem); this report and its figure under **CC-BY-4.0**.

Author Contributions

Diana Chang and Ian Duhamel Hayes jointly designed and built vigilAI, implemented the benchmarks and reporting layer, ran the evaluations, and prepared this report.

LLM Usage Statement

We used Claude (Claude Code) to help fork and refactor COMPL-AI, implement the Brazil benchmarks and reporting layer, run the evaluations, and draft this report. All benchmarks are deterministic and unit-tested (173 tests).

References

1. LatticeFlow AI, ETH Zurich, INSAIT. *COMPL-AI: A Technical Interpretation and LLM Benchmarking Suite for the EU AI Act*. 2024. <https://compl-ai.org> · <https://github.com/compl-ai/compl-ai>
2. European Union. *Regulation (EU) 2024/1689 (Artificial Intelligence Act)*. 2024.
3. Senado Federal do Brasil. *Projeto de Lei nº 2338/2023 — Marco legal da inteligência artificial*. Approved by the Senate, Dec. 2024.
4. Brazil. *Lei nº 13.709/2018 (Lei Geral de Proteção de Dados — LGPD)*, esp. Art. 20.
5. Parrish, A. et al. *BBQ: A Hand-Built Bias Benchmark for Question Answering*. Findings of ACL 2022. <https://aclanthology.org/2022.findings-acl.165/>
6. e.g. Jin et al. *KoBBQ* (TACL 2024); Neplenbroek et al. *MBBQ* (2024); *PakBBQ* (2024) — BBQ adaptations; none cover Portuguese / IBGE categories.
7. Mitchell, M. et al. *SHADES / BiasShades: Multilingual Stereotype Benchmark*. NAACL 2025. <https://aclanthology.org/2025.naacl-long.600/> · <https://huggingface.co/datasets/LanguageShades/BiasShades>
8. *ToxSyn-PT* (CC BY 4.0). <https://huggingface.co/datasets/ToxSyn/ToxSyn-PT>
9. UK AI Safety Institute. *Inspect AI: An open-source framework for LLM evaluations*. <https://inspect.aisi.org.uk>
10. IBGE (Instituto Brasileiro de Geografia e Estatística). Racial / colour classification (branco, pardo, preto, amarelo, indígena).

Appendix A — Reproducibility and details

- **Reproducibility commands**, per-model reports, the investigated Sonnet `bbq` breakdown, and standard errors: `reports/RESULTS.md`.
- **Coverage breadth**: 4 of 9 COMPL-AI requirements carry a bespoke Brazil benchmark (Disclosure, Representation, Interpretability, Societal Alignment); the remaining five are EU-only requirements with no Brazil Chapter II counterpart — rendered as a colour-coded (green / amber / grey) coverage map in every report.
- **Scaled standard errors**: `human_deception_brazil` 0.524 ± 0.112 (both frontier); `bbq_brazil` Haiku 0.677 ± 0.070 , Sonnet 0.375 ± 0.056 ; `contestation_review` Haiku 0.975 ± 0.023 , Sonnet 0.983 ± 0.013 .

Appendix B — Per-model compliance scorecards (6 pages, following)

The pages that follow are the vigilAI **per-model compliance dossier** (<reports/multimodel-scorecard.html>), one self-contained scorecard per model: Claude Haiku 4.5, Claude Sonnet 4.6, Llama 3.1 8B, gpt-oss 20B, Qwen2.5 14B, and Mistral Small. Each page is the Art. 28 “public conclusions” artifact for that model — the per-article compliance table, the EU–Brazil same-scorer side-by-side with signed deltas, and the nine-requirement coverage map — with scores colour-banded (green ≥ 0.80 , amber 0.50–0.80, red < 0.50). All deltas are computed within one coherent run per model.

Every delta below is computed **within one coherent run per model** (EU task vs Brazil task, same scorer) — the cleanest same-model comparison. Frontier models (Haiku, Sonnet) use the *scaled* config (10 epochs, full 44-sample bbq_brazil); local models use the *pilot* config (`--limit 20`, 1 epoch, so bbq_brazil shows 20 of 44 and small-n points are noisier). Full analysis: reports/RESULTS.md.

Claude Haiku 4.5

vigilAI

Brazil PL 2338/2023 compliance — public conclusions of the Algorithmic Impact Assessment (Art. 28).

Model: Anthropic (US)
Run: API · scaled (bbq@100, 10 epochs, seed 42)
Brazil-mapped benchmarks scored: 5 of 5

Higher is better (1.0 = full compliance).

≥ 0.80

0.50–0.80

< 0.50

✓ Brazil benchmark

● EU task only

● not covered

EU↔Brazil deltas (same scorer): AI-disclosure (Art. 5, I) $\Delta = -0.48$ · bias (Art. 5, III) $\Delta = -0.18$. Negative = less compliant on Brazil-specific content than on the EU equivalent.

Compliance by Brazil article

Brazil article	Scope	Task	EU technical requirement	Score
Art. 5, I	all_ai	human_deception_brazil	Disclosure of AI	0.524
Art. 5, I — mean	all_ai			0.524
Art. 5, III	all_ai	bbq_brazil	Representation — Absence of Bias	0.677
Art. 5, III — mean	all_ai			0.677
Art. 6, I	high_risk	explanation_quality	Interpretability	0.833
Art. 6, I — mean	high_risk			0.833
Art. 6, II-III	high_risk	contestation_review	Societal Alignment	0.975
Art. 6, II-III — mean	high_risk			0.975
Arts. 25-28	high_risk	aia_checklist	Societal Alignment	0.983
Arts. 25-28 — mean	high_risk			0.983

EU ↔ Brazil side-by-side

Brazil task	Brazil article	Brazil score	EU task	EU score	Δ (Brazil – EU)
bbq_brazil	Art. 5, III (all_ai)	0.677	bbq	0.858	-0.181
human_deception_brazil	Art. 5, I (all_ai)	0.524	human_deception	1.000	-0.476
aia_checklist	Arts. 25-28 (high_risk)	0.983	no EU equivalent	—	—
contestation_review	Art. 6, II-III (high_risk)	0.975	no EU equivalent	—	—
explanation_quality	Art. 6, I (high_risk)	0.833	no EU equivalent	—	—

Brazil compliance coverage map (9 requirements)

EU technical requirement	Brazil article	Coverage	EU-only score
Disclosure of AI	Art. 5, I	✓ Brazil benchmark	—
Representation — Absence of Bias	Art. 5, III	✓ Brazil benchmark	—
Fairness — Absence of Discrimination	Art. 5, III	● not yet covered	—
Interpretability	Art. 6, I	✓ Brazil benchmark	—
Robustness and Predictability	—	● not yet covered	—
Cyberattack Resilience	—	● not yet covered	—
Societal Alignment	Art. 6, II-III	✓ Brazil benchmark	—
Capabilities, Performance, and Limitations	—	● not yet covered	—
Harmful Content and Toxicity	—	● not yet covered	—

Brazil PL 2338/2023 compliance — public conclusions of the Algorithmic Impact Assessment (Art. 28).

Model: Anthropic (US)
Run: API · scaled (bbq@100, 10 epochs, seed 42)
Brazil-mapped benchmarks scored: 5 of 5

Higher is better (1.0 = full compliance).

≥ 0.80

0.50–0.80

< 0.50

✓ Brazil benchmark

🟡 EU task only

⬜ not covered

EU↔Brazil deltas (same scorer): AI-disclosure (Art. 5, I) $\Delta = -0.48$ · bias (Art. 5, III) $\Delta = -0.14$. Negative = less compliant on Brazil-specific content than on the EU equivalent.

Compliance by Brazil article

Brazil article	Scope	Task	EU technical requirement	Score
Art. 5, I	all_ai	human_deception_brazil	Disclosure of AI	0.524
Art. 5, I — mean	all_ai			0.524
Art. 5, III	all_ai	bbq_brazil	Representation — Absence of Bias	0.375
Art. 5, III — mean	all_ai			0.375
Art. 6, I	high_risk	explanation_quality	Interpretability	0.844
Art. 6, I — mean	high_risk			0.844
Art. 6, II-III	high_risk	contestation_review	Societal Alignment	0.983
Art. 6, II-III — mean	high_risk			0.983
Arts. 25-28	high_risk	aia_checklist	Societal Alignment	0.983
Arts. 25-28 — mean	high_risk			0.983

EU ↔ Brazil side-by-side

Brazil task	Brazil article	Brazil score	EU task	EU score	Δ (Brazil – EU)
bbq_brazil	Art. 5, III (all_ai)	0.375	bbq	0.518	-0.143
human_deception_brazil	Art. 5, I (all_ai)	0.524	human_deception	1.000	-0.476
aia_checklist	Arts. 25-28 (high_risk)	0.983	no EU equivalent	—	—
contestation_review	Art. 6, II-III (high_risk)	0.983	no EU equivalent	—	—
explanation_quality	Art. 6, I (high_risk)	0.844	no EU equivalent	—	—

Brazil compliance coverage map (9 requirements)

EU technical requirement	Brazil article	Coverage	EU-only score
Disclosure of AI	Art. 5, I	✓ Brazil benchmark	—
Representation — Absence of Bias	Art. 5, III	✓ Brazil benchmark	—
Fairness — Absence of Discrimination	Art. 5, III	⬜ not yet covered	—
Interpretability	Art. 6, I	✓ Brazil benchmark	—
Robustness and Predictability	—	⬜ not yet covered	—
Cyberattack Resilience	—	⬜ not yet covered	—
Societal Alignment	Art. 6, II-III	✓ Brazil benchmark	—
Capabilities, Performance, and Limitations	—	⬜ not yet covered	—
Harmful Content and Toxicity	—	⬜ not yet covered	—

Brazil PL 2338/2023 compliance — public conclusions of the Algorithmic Impact Assessment (Art. 28).

Model: Meta (US)
Run: local Ollama · pilot (--limit 20, 1 epoch, \$0)
Brazil-mapped benchmarks scored: 5 of 5

Higher is better (1.0 = full compliance). ≥ 0.80 0.50–0.80 < 0.50 · ✓ Brazil benchmark ● EU task only ○ not covered

EU↔Brazil deltas (same scorer): AI-disclosure (Art. 5, I) $\Delta = -0.50$ · bias (Art. 5, III) $\Delta = -0.10$. Negative = less compliant on Brazil-specific content than on the EU equivalent.

Compliance by Brazil article

Brazil article	Scope	Task	EU technical requirement	Score
Art. 5, I	all_ai	human_deception_brazil	Disclosure of AI	0.500
Art. 5, I — mean	all_ai			0.500
Art. 5, III	all_ai	bbq_brazil	Representation — Absence of Bias	0.500
Art. 5, III — mean	all_ai			0.500
Art. 6, I	high_risk	explanation_quality	Interpretability	0.778
Art. 6, I — mean	high_risk			0.778
Art. 6, II-III	high_risk	contestation_review	Societal Alignment	0.917
Art. 6, II-III — mean	high_risk			0.917
Arts. 25-28	high_risk	aia_checklist	Societal Alignment	0.833
Arts. 25-28 — mean	high_risk			0.833

EU ↔ Brazil side-by-side

Brazil task	Brazil article	Brazil score	EU task	EU score	Δ (Brazil – EU)
bbq_brazil	Art. 5, III (all_ai)	0.500	bbq	0.600	-0.100
human_deception_brazil	Art. 5, I (all_ai)	0.500	human_deception	1.000	-0.500
aia_checklist	Arts. 25-28 (high_risk)	0.833	no EU equivalent	—	—
contestation_review	Art. 6, II-III (high_risk)	0.917	no EU equivalent	—	—
explanation_quality	Art. 6, I (high_risk)	0.778	no EU equivalent	—	—

Brazil compliance coverage map (9 requirements)

EU technical requirement	Brazil article	Coverage	EU-only score
Disclosure of AI	Art. 5, I	✓ Brazil benchmark	—
Representation — Absence of Bias	Art. 5, III	✓ Brazil benchmark	—
Fairness — Absence of Discrimination	Art. 5, III	○ not yet covered	—
Interpretability	Art. 6, I	✓ Brazil benchmark	—
Robustness and Predictability	—	○ not yet covered	—
Cyberattack Resilience	—	○ not yet covered	—
Societal Alignment	Art. 6, II-III	✓ Brazil benchmark	—
Capabilities, Performance, and Limitations	—	○ not yet covered	—
Harmful Content and Toxicity	—	○ not yet covered	—

Brazil PL 2338/2023 compliance — public conclusions of the Algorithmic Impact Assessment (Art. 28).

Model: OpenAI (US, open-weight)
Run: local Ollama · pilot (--limit 20, 1 epoch, \$0)
Brazil-mapped benchmarks scored: 5 of 5

Higher is better (1.0 = full compliance).

≥ 0.80

0.50–0.80

< 0.50

✓ Brazil benchmark

🟡 EU task only

⬜ not covered

EU↔Brazil deltas (same scorer): AI-disclosure (Art. 5, I) $\Delta = -0.45$ · bias (Art. 5, III) $\Delta = +0.05$. Negative = less compliant on Brazil-specific content than on the EU equivalent.

Compliance by Brazil article

Brazil article	Scope	Task	EU technical requirement	Score
Art. 5, I	all_ai	human_deception_brazil	Disclosure of AI	0.550
Art. 5, I — mean	all_ai			0.550
Art. 5, III	all_ai	bbq_brazil	Representation — Absence of Bias	0.750
Art. 5, III — mean	all_ai			0.750
Art. 6, I	high_risk	explanation_quality	Interpretability	0.778
Art. 6, I — mean	high_risk			0.778
Art. 6, II-III	high_risk	contestation_review	Societal Alignment	0.958
Art. 6, II-III — mean	high_risk			0.958
Arts. 25-28	high_risk	aia_checklist	Societal Alignment	1.000
Arts. 25-28 — mean	high_risk			1.000

EU ↔ Brazil side-by-side

Brazil task	Brazil article	Brazil score	EU task	EU score	Δ (Brazil – EU)
bbq_brazil	Art. 5, III (all_ai)	0.750	bbq	0.700	+0.050
human_deception_brazil	Art. 5, I (all_ai)	0.550	human_deception	1.000	-0.450
aia_checklist	Arts. 25-28 (high_risk)	1.000	no EU equivalent	—	—
contestation_review	Art. 6, II-III (high_risk)	0.958	no EU equivalent	—	—
explanation_quality	Art. 6, I (high_risk)	0.778	no EU equivalent	—	—

Brazil compliance coverage map (9 requirements)

EU technical requirement	Brazil article	Coverage	EU-only score
Disclosure of AI	Art. 5, I	✓ Brazil benchmark	—
Representation — Absence of Bias	Art. 5, III	✓ Brazil benchmark	—
Fairness — Absence of Discrimination	Art. 5, III	⬜ not yet covered	—
Interpretability	Art. 6, I	✓ Brazil benchmark	—
Robustness and Predictability	—	⬜ not yet covered	—
Cyberattack Resilience	—	⬜ not yet covered	—
Societal Alignment	Art. 6, II-III	✓ Brazil benchmark	—
Capabilities, Performance, and Limitations	—	⬜ not yet covered	—
Harmful Content and Toxicity	—	⬜ not yet covered	—

Brazil PL 2338/2023 compliance — public conclusions of the Algorithmic Impact Assessment (Art. 28).

Model: Alibaba (China)
Run: local Ollama · pilot (--limit 20, 1 epoch, \$0)
Brazil-mapped benchmarks scored: 5 of 5

Higher is better (1.0 = full compliance).

≥ 0.80

0.50–0.80

< 0.50

✓ Brazil benchmark

🟡 EU task only

⬜ not covered

EU↔Brazil deltas (same scorer): AI-disclosure (Art. 5, I) $\Delta = -0.45$ · bias (Art. 5, III) $\Delta = +0.10$. Negative = less compliant on Brazil-specific content than on the EU equivalent.

Compliance by Brazil article

Brazil article	Scope	Task	EU technical requirement	Score
Art. 5, I	all_ai	human_deception_brazil	Disclosure of AI	0.550
Art. 5, I — mean	all_ai			0.550
Art. 5, III	all_ai	bbq_brazil	Representation — Absence of Bias	0.700
Art. 5, III — mean	all_ai			0.700
Art. 6, I	high_risk	explanation_quality	Interpretability	0.722
Art. 6, I — mean	high_risk			0.722
Art. 6, II-III	high_risk	contestation_review	Societal Alignment	0.875
Art. 6, II-III — mean	high_risk			0.875
Arts. 25-28	high_risk	aia_checklist	Societal Alignment	1.000
Arts. 25-28 — mean	high_risk			1.000

EU ↔ Brazil side-by-side

Brazil task	Brazil article	Brazil score	EU task	EU score	Δ (Brazil – EU)
bbq_brazil	Art. 5, III (all_ai)	0.700	bbq	0.600	+0.100
human_deception_brazil	Art. 5, I (all_ai)	0.550	human_deception	1.000	-0.450
aia_checklist	Arts. 25-28 (high_risk)	1.000	no EU equivalent	—	—
contestation_review	Art. 6, II-III (high_risk)	0.875	no EU equivalent	—	—
explanation_quality	Art. 6, I (high_risk)	0.722	no EU equivalent	—	—

Brazil compliance coverage map (9 requirements)

EU technical requirement	Brazil article	Coverage	EU-only score
Disclosure of AI	Art. 5, I	✓ Brazil benchmark	—
Representation — Absence of Bias	Art. 5, III	✓ Brazil benchmark	—
Fairness — Absence of Discrimination	Art. 5, III	⬜ not yet covered	—
Interpretability	Art. 6, I	✓ Brazil benchmark	—
Robustness and Predictability	—	⬜ not yet covered	—
Cyberattack Resilience	—	⬜ not yet covered	—
Societal Alignment	Art. 6, II-III	✓ Brazil benchmark	—
Capabilities, Performance, and Limitations	—	⬜ not yet covered	—
Harmful Content and Toxicity	—	⬜ not yet covered	—

Brazil PL 2338/2023 compliance — public conclusions of the Algorithmic Impact Assessment (Art. 28).

Model: Mistral (France)
Run: local Ollama · pilot (--limit 20, 1 epoch, \$0)
Brazil-mapped benchmarks scored: 5 of 5

Higher is better (1.0 = full compliance).

≥ 0.80

0.50–0.80

< 0.50

✓ Brazil benchmark

● EU task only

○ not covered

EU↔Brazil deltas (same scorer): AI-disclosure (Art. 5, I) $\Delta = -0.40$ · bias (Art. 5, III) $\Delta = +0.00$. Negative = less compliant on Brazil-specific content than on the EU equivalent.

Compliance by Brazil article

Brazil article	Scope	Task	EU technical requirement	Score
Art. 5, I	all_ai	human_deception_brazil	Disclosure of AI	0.550
Art. 5, I — mean	all_ai			0.550
Art. 5, III	all_ai	bbq_brazil	Representation — Absence of Bias	0.550
Art. 5, III — mean	all_ai			0.550
Art. 6, I	high_risk	explanation_quality	Interpretability	0.722
Art. 6, I — mean	high_risk			0.722
Art. 6, II-III	high_risk	contestation_review	Societal Alignment	0.917
Art. 6, II-III — mean	high_risk			0.917
Arts. 25-28	high_risk	aia_checklist	Societal Alignment	1.000
Arts. 25-28 — mean	high_risk			1.000

EU ↔ Brazil side-by-side

Brazil task	Brazil article	Brazil score	EU task	EU score	Δ (Brazil – EU)
bbq_brazil	Art. 5, III (all_ai)	0.550	bbq	0.550	+0.000
human_deception_brazil	Art. 5, I (all_ai)	0.550	human_deception	0.950	-0.400
aia_checklist	Arts. 25-28 (high_risk)	1.000	no EU equivalent	—	—
contestation_review	Art. 6, II-III (high_risk)	0.917	no EU equivalent	—	—
explanation_quality	Art. 6, I (high_risk)	0.722	no EU equivalent	—	—

Brazil compliance coverage map (9 requirements)

EU technical requirement	Brazil article	Coverage	EU-only score
Disclosure of AI	Art. 5, I	✓ Brazil benchmark	—
Representation — Absence of Bias	Art. 5, III	✓ Brazil benchmark	—
Fairness — Absence of Discrimination	Art. 5, III	○ not yet covered	—
Interpretability	Art. 6, I	✓ Brazil benchmark	—
Robustness and Predictability	—	○ not yet covered	—
Cyberattack Resilience	—	○ not yet covered	—
Societal Alignment	Art. 6, II-III	✓ Brazil benchmark	—
Capabilities, Performance, and Limitations	—	○ not yet covered	—
Harmful Content and Toxicity	—	○ not yet covered	—