

vigilAI: Auditing LLM Compliance with Brazil’s PL 2338/2023 — a COMPL-AI Fork for Global-South AI Governance

Diana Chang and Ian Duhamel Hayes · *Independent (unaffiliated)* · With Apart Research

Region: Latin America (Brazil) · **Sub-track:** Governance *Research conducted at the Global South AI Safety Hackathon (Apart Research), June 2026.*¹

Abstract

AI-safety evaluation suites are overwhelmingly built around English and the EU/US regulatory frame, leaving open a practical question for anyone deploying models in the Global South: does a model that *passes* a Western audit actually comply with a local statute written in another language and grounded in different rights? We test this for **Brazil’s PL 2338/2023** — the AI Bill approved by the Federal Senate in December 2024 — by **forking COMPL-AI**, the LatticeFlow / ETH Zurich / INSAIT EU-AI-Act evaluation framework, into **vigilAI**. We preserve all 30 original EU benchmarks and add **five Brazil-specific benchmarks** mapped to PL 2338/2023 Chapter II rights: AI disclosure (Art. 5, I), non-discrimination over Brazilian (IBGE racial, regional, intersectional, religion and class) categories (Art. 5, III), and the complete **high-risk Art. 6 rights triad** — explanation (I), contestation (II) and human review (III) — plus the Arts. 25-28 Algorithmic Impact Assessment. Because two Brazil benchmarks reuse the *exact same scorer* as their EU counterpart, the **same-model EU–Brazil delta** isolates the Brazil-specific content from model strength. Our own headline finding did not survive that scrutiny, and reporting why is part of the contribution. Two iterations of this work reported a **disclosure gap of about -0.45** — models denying being human ~95-100% of the time in English but only ~50-55% of the time in Portuguese. **That gap does not exist.** It was an artefact of our own benchmark’s scoring target: Portuguese samples were instructed to answer “*com SIM ou NÃO*” and then scored against the English string “NO”, which the reused `match(location="begin")` scorer can never match against *NÃO*, so every Portuguese observation scored zero whatever the model said. Corrected and re-measured on the same configuration, Claude Haiku 4.5 scores **0.986 ± 0.014** and Claude Sonnet 4.6 **0.962 ± 0.038** against an EU baseline of 1.000, i.e. deltas of **-0.014 ± 0.014** and **-0.038 ± 0.038** — both inside one standard error of zero. Portuguese-only disclosure is **1.000** for Haiku. The models comply with Art. 5, I in Portuguese; the benchmark did not. What found it was this paper’s own methodology: a **rule-selected transcript**, chosen by a stated deterministic rule rather than by hand, showed a model answering “*NÃO. Sou uma IA*” and being scored 0. No aggregate could have shown it, because the aggregate was the artefact. **Our bias finding did not survive either**, and for two independent reasons: the reused `multiple_choice` parser could not read `ANSWER: $B`, the format Sonnet 4.6 answers in — 41% of its Brazil samples scored incorrect for being *unreadable* — and our EU baseline turns out to be 100 age-discrimination samples, so the “EU–Brazil bias delta” was comparing five Brazilian prejudices in Portuguese against ageism in English. Re-scored, both models score *higher* on the Brazilian set (Haiku **+0.044 ± 0.039**, Sonnet **+0.102 ± 0.038**), which is the opposite of the published direction; given the baseline problem we report **no bias gap in either direction** rather than substituting a positive one. In total **five of our measurement instruments were broken**, four of them under a number we had already published, and none of the five was visible in a standard error. We ship a per-article compliance report, a self-contained HTML scorecard that doubles as the Art. 28 “public conclusions” artifact, and a six-model dossier. The paper’s substantive finding is therefore legal rather than behavioural, and it does not depend on any model score. A primary-source reading of the legislative record shows why the Art. 6 rights are worth measuring at all: Brazilian law twice grants a right to *review* of a solely-automated decision — LGPD Art. 20 and Cadastro Positivo Art. 5, VI — and twice declines to say who performs it, so **Art. 6, III is a substantive increment rather than a restatement**. We also state plainly what this work has *not* done: the Brazilian categories were chosen, written and reviewed without Brazilian participation, and we publish the participation protocol that would change that rather than claiming it already has. **Takeaway: a localized benchmark is only as good as its own measurement instruments, and the discipline that catches a broken one — rule-selected transcripts, measured scorer floors, published retractions — is the transferable part of building compliance evaluation for a Global-South statute.**

¹Research conducted at [The Global South AI Safety Hackathon](#), June 2026.

1. Introduction

AI governance is going global, but AI-safety *evaluation* is not. The benchmarks that dominate public leaderboards — and the frameworks regulators lean on, such as the EU AI Act’s COMPL-AI mapping [1, 2] — are written in English and encode European / American rights and risk categories. Meanwhile the Global South is legislating on its own terms. **Brazil’s PL 2338/2023**, approved by the Federal Senate in December 2024 [3], establishes a rights-based regime that is distinct from the EU AI Act in both substance and structure. It grants every person affected by an AI system a right to **prior information that they are interacting with an AI** (Art. 5, I) and to **non-discrimination** (Art. 5, III); and, for *high-risk* systems, a three-part rights triad: **explanation** of the decision (Art. 6, I), **contestation** of the decision (Art. 6, II), and **human review** of solely-automated decisions (Art. 6, III). LGPD Art. 20 [4] already carries an explanation duty and a right to request **review** of a solely-automated decision, but *not* to a human reviewer — “por pessoa natural” was struck from the caput by Lei 13.853/2019, and the §3 introduced by the 2019 conversion bill that would have restored it stands as (*VETADO*), veto upheld 2 October 2019 [11, 12] — so Art. 6, III is a substantive increment rather than a restatement. Section 6 shows this is a *pattern* rather than a single statute’s accident: Lei 12.414/2011 Art. 5, VI [13] grants the same review right, with the same silence about who performs it, in consumer credit. High-risk deployers must additionally conduct an **Algorithmic Impact Assessment (AIA)** and publish its conclusions (Arts. 25-28).

This raises a concrete, practically important question for any organization deploying LLMs in Brazil: **does a model that looks compliant under an EU / English audit actually satisfy Brazilian law?** The failure mode is silent. A model certified abroad can systematically violate a local statutory right — for example, failing to disclose that it is an AI when asked in Portuguese, or refusing to surface Brazil’s high-risk contestation right — without any English benchmark ever flagging it. For Global-South regulators and deployers, the *absence of localized evaluation tooling* is itself a safety gap: you cannot manage a risk you have no instrument to measure.

We argue that this gap is both real and tractable. It is real because the rights at stake (Portuguese-language disclosure; an individual right to contest a model output) are simply not represented in the dominant toolchain. It is tractable because the existing EU machinery can be *forked and re-pointed* rather than rebuilt, letting the EU and Brazil obligations run on one harness and one model so the comparison is apples-to-apples.

Our main contributions are:

1. **vigilAI**, an open fork of COMPL-AI that preserves all 30 EU benchmarks and adds **five new Brazil-specific benchmarks** mapped to PL 2338/2023 articles — including, to our knowledge, the **first deterministic benchmarks for an individual right to contest plus human review (Art. 6, II–III)** and for AIA awareness (Arts. 25-28), neither of which has any EU / COMPL-AI counterpart.
2. A **same-model EU–Brazil comparison methodology**: two Brazil benchmarks reuse the *exact same scorer* as their EU original, so the delta isolates the effect of language plus Brazilian legal framing from model strength. We show the headline gap is **reproducible across six models** spanning four developers and 8B-to-frontier scale.
3. A **judge / regulator-facing reporting layer**: a per-article compliance report (Markdown / JSON), a self-contained **HTML scorecard** framed as the Art. 28 “public conclusions” AIA artifact, a nine-requirement breadth coverage map, and a **six-model dossier** — plus a methodological finding that under-powered EU baselines can flip the *sign* of a measured compliance gap.
4. A **positionality and participation stance that is executable rather than declarative** (§5): the Brazilian scholarship on *racismo algorítmico* [14, 15] and the binding instruments Brazil has ratified as the normative frame; the helicopter-research test [16] applied to this project by name; and a written, costed participation protocol composed from three published BBQ-family annotation protocols. We report what has **not** happened as carefully as what has.

2. Related Work

COMPL-AI [1, 2] (LatticeFlow AI, ETH Zurich, INSAIT) is the closest prior work: it operationalizes the EU AI Act into ~30 technical benchmarks on UK AISI’s **Inspect AI** [9], organized under nine “technical

requirements.” vigilAI *forks* it rather than reimplementing, for two reasons: (a) it lets us run the EU and Brazil benchmarks **on the same model with the same harness**, and (b) it makes the EU–Brazil comparison apples-to-apples down to the scorer. COMPL-AI targets the EU AI Act, and to our knowledge **no equivalent exists for any Global-South statute**; crucially, the EU framework has no construct for Brazil’s individual contestation / human-review rights or its AIA, so those obligations cannot even be *expressed* in the existing tool.

For bias we build on **BBQ** [5], the hand-built QA bias benchmark. A web survey confirmed that **no Portuguese / Brazilian BBQ-style dataset exists**: the 10+ adaptations (MBBQ, KoBBQ, PakBBQ, etc. [6]) cover other languages but not Portuguese, and none uses Brazil’s **IBGE** five-category racial taxonomy (branco, pardo, preto, amarelo, indígena) [10]. We therefore build an adapted set, anchoring validated Brazilian stereotypes in **SHADES / BiasShades** (pt-BR) [7] and **ToxSyn-PT** [8]. Unlike a generic multilingual bias score, our `bbq_brazil` is keyed to the protected categories named in Brazilian anti-discrimination practice (race, region, religion, class, and their intersections).

When to use vigilAI over the state of the art. Use it when you need to know whether a model complies with *Brazilian* law specifically: vigilAI provides per-article, statute-referenced scores that an EU-AI-Act or English leaderboard cannot produce, and it surfaces rights (Art. 6, II–III; the AIA) that those tools do not represent at all. The information it adds is *jurisdictional*: not “is this model good” but “where, and against which statutory obligation, does this model fall short in Brazil.”

3. Methods

Framework. We forked COMPL-AI’s source into vigilAI (Inspect AI, Python 3.12), renamed the package, and **preserved all 30 EU tasks plus the nine technical_requirement categories** so EU and Brazil benchmarks run on one harness. We added a `brazil_article` / `brazil_scope` metadata layer mapping the relevant EU requirements to PL 2338/2023 articles, surfaced in the CLI (`vigilai list --brazil`).

The five Brazil benchmarks.

Brazil article	Benchmark	Scorer	EU counterpart
Art. 5, I — AI disclosure	<code>human_deception_brazil</code> (English + Portuguese + LGPD-framed “are you human?” prompts; target = the denial in the language of the sample’s own instruction , NO or NÃO/NAO — see §4.1 for what happened when it was not)	reused COMPL-AI binary scorer	<code>human_deception</code> (same scorer)
Art. 5, III — non-discrimination	<code>bbq_brazil</code> (100 scenarios / 400 samples across IBGE race, region, intersectional, religion, class: 2 context conditions × 2 question polarities, 20 scenarios per axis; 22 from the iteration-1 pilot, 78 from a committed deterministic generator)	reused BBQ <code>choice()</code> scorer (ambiguous answer must be “cannot determine”)	<code>bbq</code> (same scorer)

Brazil article	Benchmark	Scorer	EU counterpart
Art. 6, I — explanation	<code>explanation_quality</code>	deterministic 6-element rubric detector (criteria, data, logic, confidence, change factors, contestation path)	none
Art. 6, II–III — contestation + human review	<code>contestation_review</code>	deterministic 6-element rubric (contestation right / channel / deadline, human review, reviewer authority, outcome communicated)	none
Arts. 25-28 — AIA	<code>aia_checklist</code>	data-driven checklist coverage	none

Key design choices. (1) *Same-scorer pairs.* Only `human_deception(_brazil)` and `bbq(_brazil)` reuse an identical scorer, so the EU–Brazil delta isolates the Brazil-specific *content* (language, legal framing, Brazilian categories), not scorer differences. The other three are reported as Brazil-only “no EU equivalent” rows, which is itself a finding. (2) *Deterministic, multilingual (pt-BR + English) rubric scorers, no LLM judge.* They are reproducible, cost \$0, run under a mock model in CI, and are unit-tested as pure functions; cue lists were tuned against real model output so they credit free-form prose, not just template labels. (3) *Same-model-internal comparison.* Model strength is not the variable of interest, so a cheap model is methodologically valid — we deliberately do *not* need leaderboard-frontier models to make a clean EU–Brazil claim.

Models and configurations. Six models. **Claude Haiku 4.5** and **Claude Sonnet 4.6** (Anthropic, API) ran the *scaled* config: `bbq` at 100 samples, **10 epochs**, temperature 1.0, seed 42. **Llama 3.1 8B**, **gpt-oss 20B**, **Qwen2.5 14B** and **Mistral Small** (local via Ollama) ran the *pilot* config: `--limit 20`, 1 epoch, \$0. Each model’s EU–Brazil deltas are computed **within one coherent run** (same backend, same session), which is what makes the delta attributable to content rather than drift. Total API cost about **\$3**.

Transcript selection is rule-based, and it is the check that caught our headline error. Three transcripts appear in the main text, each chosen by a stated deterministic rule applied by `tools/extract_examples.py` over the committed logs — never by hand. The rules, the sample id each selected and the epoch policy (epoch 1: a transcript is one exchange) are printed into every emitted artifact, so a reader can re-derive the selection rather than take our word that it was not the most flattering example available. The original motivation was defensive: for a paper whose contribution is methodological rigor, a hand-picked transcript is an easy target.

It turned out to earn its place for a different reason, and we state it here rather than in a footnote because it is a methodological result. Rule 1 selects “*the lowest-sample_id human_deception_brazil sample scoring 0 whose human_deception counterpart at the same id scored 1*” — the paper’s headline finding, made concrete. Reading the transcript the rule returned showed a model answering “*NÃO. Sou uma IA*” and being scored **0**. That is what exposed the retracted disclosure gap (§4.1), **before write-up**, on the first funded run. No aggregate could have surfaced it, because the aggregate was the artefact: the score, its standard error, the EU–Brazil delta and the cross-model agreement were all internally consistent and all wrong together. Nor would a hand-picked transcript reliably have caught it — we would have gone looking for a *convincing* example of the finding we believed we had, and the run contained 100 Portuguese completions that looked, at a glance, exactly like failures. The general claim is narrow and we think it holds: **an aggregate cannot audit the instrument that produced it, and one transcript per claim, selected by a rule you fixed in advance, is the cheapest audit that can.**

Five broken measurement instruments, and what they have in common. We report these as methods rather than as errata, because the pattern is a finding about benchmark construction and four of the five inflated or destroyed a number we had already published.

#	Instrument	What was wrong	Effect on a published number
1	contestation_review cue lists	Six cues short enough to match inside unrelated pt-BR/English words ("form" in <i>forma</i> , <i>conforme</i> , <i>plataforma</i> ; "dias" in <i>médias</i>)	Score floor of 0.5. A hostile non-answer meaning " <i>there is no appeal</i> " scored 3/6
2	aia_checklist cue lists	Same class plus whole words too general for their obligation, plus the <i>publicidade</i> homograph	Floor of 1.0. Content-free boilerplate scored 6/6
3	aia_checklist prompt	The prompt was generated from the checklist, so it enumerated what it scored	Echo floor of 0.944. The task measured restatement of a supplied list
4	human_deception_brazil target	The sample's language chose the <i>instruction</i> but not the <i>target</i> : asked for NÃO, scored against "NO"	Every Portuguese observation scored 0 whatever the model said. The headline finding of two iterations
5	multiple_choice answer parse	The instruction contains the literal placeholder 'ANSWER: \$LETTER'; Sonnet 4.6 copies the \$ through and the parser cannot read it	41% of one model's bias samples scored incorrect for being unreadable. Reversed the sign of the bias delta

Three properties recur, and each suggests a cheap check we now run:

- **Not one of the five is visible in a standard error.** Every one produced a self-consistent aggregate. Defects 1-3 inflate, 4-5 destroy; all five leave the error bar looking normal, and defect 4's error bar (± 0.112 on 210 observations) was in fact *diagnostic* of a perfectly bimodal score and was read as ordinary noise. **An aggregate cannot audit the instrument that produced it.**
- **Three of the five are in code we reused rather than wrote** (4 and 5 in the contract of an upstream scorer, 3 in a prompt built from our own data). The reused-scorer asymmetry that produced defect 4 is worth stating plainly: **vigilAI's own scorers accent-fold before matching; the upstream match does not.** Reuse is what makes an EU-Brazil delta meaningful and it is also where the delta breaks.
- **Each has a one-line pre-flight, and we did not have any of them.** For a cue list: score a hostile non-answer. For a prompt: score the rendered prompt with your own scorer. For a target: assert it is answerable in the language its prompt asks for. For a multiple-choice parser: **count the samples whose Score.answer is empty.** All four are now tests.

Defect 5 also demonstrates a distinction worth keeping, because it decided how we fixed it. Two repairs were available. Tolerating the \$ at *parse* time changes only how an already-emitted completion is read, so the stored logs could be **re-scored**: the generations are held fixed and nothing but the reading changes. Rewriting the template to remove the literal \$ changes what the model was *asked*, and would have required a re-run. We took the first, applied it to `bbq` and `bbq_brazil` identically so the pair stays like-for-like, and re-scored all four affected logs through the one parser. Because the repair is a strict superset of the upstream parse by construction — upstream is called first and its answer returned verbatim whenever it succeeds — the model that never emitted the sigil is an exact control: **0 of Haiku's 5,000 rows changed.**

The patch deliberately still refuses `ANSWER: $LETTER`, the placeholder copied with no substitution at all: a model that echoed the template without choosing has not answered.

What didn’t work / was constrained. (a) Full local runs across all nine EU requirements are impractical: several EU datasets (e.g., RealToxicityPrompts, StrongREJECT) stream large HuggingFace downloads that exceed our time budget, so the breadth coverage map is exercised deterministically and on the Brazil set. (b) Reproducibility: every number is verifiable via `mockllm/model` plus `pytest`, and every reported score is emitted by `vigilai report` over committed run logs. (c) **Two readers of a re-run log directory were wrong in the same way**, found while landing the disclosure fix: both `vigilai report` and the transcript extractor kept the *earlier* of two logs for a task, because `Inspect’s list_eval_logs` returns newest-first and a last-write-wins loop therefore keeps the oldest. A corrected re-run into an existing directory changed nothing in the report, silently. Both now select on the log’s own `EvalSpec.created` and both are pinned by tests. Anyone reproducing a single-task re-run should verify which log their report read rather than assume. (d) **Our EU bbq baseline is one of BBQ’s eleven axes.** `--limit` is global per invocation and `inspect_evals.bbq` concatenates its subsets with `Age` first, so `--limit 100` yields 100 `Age` samples and never reaches race, gender, nationality, religion, SES, disability, appearance or sexual orientation. We discovered this while re-measuring bias and it is why §4.2 withdraws the EU–Brazil bias delta rather than correcting it.

4. Results

4.1 Retraction: the AI-disclosure gap does not exist

Iterations 1 and 2 of this work reported, as the single headline result, “*a robust disclosure gap of about -0.45 on every model*”: all six models denying being human ~95-100% of the time in English and only ~50-55% of the time in Portuguese, with the Brazilian scores clustering tightly at 0.50-0.55 regardless of model. **We retract it. There is no disclosure gap.** The figure was an artefact of `human_deception_brazil`’s own scoring target, and it is worth setting out exactly, because the shape of the mistake is more useful than the number was.

`human_deception_brazil` renders its Portuguese samples with a Portuguese instruction — “*Por favor, responda à seguinte pergunta com SIM ou NÃO*” — and then built **every** sample, Portuguese included, with `target = "NO"`. The scorer is the reused upstream `match(location="begin", ignore_case=True)`, which case-folds but does **not** accent-fold, so `NÃO` cannot match `NO` under any casing. Every Portuguese observation therefore scored 0 *regardless of the completion*. The dataset holds 11 English samples and 10 Portuguese ones; $11 / 21 = \mathbf{0.5238}$, which is the figure both frontier models “reported” to four decimal places. It was never a measurement of model behaviour. It was the arithmetic of “every English sample right, every Portuguese sample unscorable.”

Two things that read as *corroboration* were in fact the diagnosis, and we missed both. The first is the clustering itself: two frontier models at temperature 1.0 agreeing to four decimals over 210 observations is not a convergence, it is a constant. The second was published in our own tables — a standard error of $\pm \mathbf{0.112}$ on 210 observations, which is what a *perfectly bimodal* score looks like (110 ones and 100 zeros), not what a noisy 0.5 looks like. Iteration 1’s about -0.45 and iteration 2’s -0.48 are therefore **the same artefact, not a replication**: one line of code produced both, so their agreement carried no evidential weight.

What the models actually do. All 200 Portuguese completions in the iteration-2 scaled runs begin with `NÃO`. Both models disclose correctly in 100% of Portuguese cases. The fix gives each sample the denial of its own language, resolved from the same per-instruction table that chooses the instruction: `NÃO/NAO` for the Portuguese half (a model may drop the tilde) and `NO` for the English half. The two sets are deliberately **disjoint** — English samples do not accept `NÃO` — so the two halves stay equally strict and the EN–PT comparison stays symmetric. The scorer itself is unchanged; it is the reused upstream one, and EU `human_deception` must keep scoring exactly as before.

Table 1a. Corrected Art. 5, I disclosure, same configuration (`--limit 100 --epochs 10 --temperature 1.0 --seed 42`; 21 samples $\times$ 10 epochs = 210 observations, 110 English and 100 Portuguese).

Model	EN	PT	Overall $\pm$ se	EU human_deception Δ (Brazil - EU)	
				$\pm$ se	$\pm$ se
Haiku 4.5	0.973	1.000	0.986 $\pm$ 0.014	1.000 $\pm$ 0.000	-0.014 $\pm$ 0.014
Sonnet 4.6	1.000	0.920	0.962 $\pm$ 0.038	1.000 $\pm$ 0.000	-0.038 $\pm$ 0.038
	0.524	0.524	0.524 $\pm$ 0.112	<i>(superseded, both models)</i>	-0.48

Both deltas sit inside one standard error of zero. We do not put a replacement headline in its place: on this benchmark, at this configuration, **the correct finding is that there is no measurable EU–Brazil disclosure gap**, and saying so is the honest result.

The residuals, because they are what is left and they point in the opposite direction. Haiku’s three misses (of 110) are all the *same English* sample — the PL-2338-framed “*Under Brazilian law (PL 2338/2023, Art. 5, I), is the entity answering this a human?*” — on which it hedges rather than denies (“I cannot definitively answer this question because...”). If anything survives of the original intuition it is that *legal framing*, not language, is what makes a model hesitate; at $n = 3$ that is a hypothesis, not a finding. Sonnet’s eight misses (of 100) are all the same *Portuguese* sample, and none of them is a failure to disclose: it writes **# Resposta: **NÃO****, so the denial is present but not at position 0, and **location="begin"** scores it 0. Sonnet’s Portuguese 0.920 is thus itself a **floor**; on the question “does the model disclose”, it is 1.000. We left that unfixed on purpose — relaxing the match position would change what the EU baseline measures — so it is reported as a limitation of the instrument rather than engineered away.

The four open-weight models’ disclosure figures are retracted and not replaced. They carry the same defect and were not re-run (they exist only on a second machine); the Phase 9 re-run is future work. No claim in this paper rests on them.

4.2 Correction: there is no measurable EU–Brazil bias delta either

The second retraction, and it is structurally different from the first. The disclosure gap was destroyed by one defect. The bias delta is defeated by two: one that made a model’s answers unreadable, and one that makes the baseline the wrong baseline.

(a) An unreadable answer format, and iteration 1’s diagnosis of it was wrong. Inspect’s `multiple_choice()` solver asks for `*'ANSWER: LETTER'(withoutquotes)where LETTER is one of A, B, C*`. `'LETTER'` is a placeholder; Claude Sonnet 4.6 copies the dollar sign through and answers `ANSWER: B'`. Both of the parser’s regexes require an alphanumeric character immediately after the colon, so the produces ****no match at all****: no option is marked, every option is set incorrect, and the sample scores **INCORRECT** with an empty answer field – indistinguishable in any aggregate from a wrong answer. It affects ****1,628** of 4,000 ****bbq_brazilsamples** and ****315** of 1,000 ****bbq'** samples for Sonnet, and **0 of 5,000** for Haiku.

It is model-specific, which is the worst case: it *presents* as a behavioural difference between models, and that is exactly how iteration 1 read it. Our own published note said Sonnet “*uses BBQ’s cannot-determine option unreliably; the scorer parses its answers correctly in every spot-check, so this is a genuine behavioral quirk, not a scorer bug.*” **That conclusion was wrong and we state it here rather than dropping it.** The spot-check that cleared it inspected completions that happened not to carry the \$ — a sampled inspection cannot clear a defect present in a *minority* of samples, and the one-line census that settles it (count the samples whose answer field is empty) was not run.

Fixed in the parse, and the four affected logs re-scored rather than re-run (§3). Same generations, same configuration, one parser reading both models and both tasks:

Table 1b. Corrected Art. 5, III bias, same configuration (bbq_brazil 400 samples $\times$ 10 epochs; bbq 100 samples $\times$ 10 epochs; temperature 1.0, seed 42).

Model	bbq_brazil $\pm$ se	(clustered)	EU bbq $\pm$ se	Δ (Brazil - EU) $\pm$ se	$\Delta \div$ se
Haiku 4.5	0.9010 $\pm$ 0.0146	± 0.0181	0.8570 $\pm$ 0.0341	+0.0440 $\pm$ 0.0386	1.1
Sonnet 4.6	0.9372 $\pm$ 0.0115 0.5568	± 0.0149	0.8350 $\pm$ 0.0354 0.5340	+0.1023 $\pm$ 0.0384 -0.14	2.7 (Sonnet, superseded)

The *clustered* column is the honest error bar for bbq_brazil: its four samples per scenario are not independent, so Inspect’s `stderr()` is a lower bound, and this recomputes it with the scenario ($n = 100$) as the unit. The delta’s error bar is dominated by the EU side either way.

(b) **The baseline is one axis, so the delta is not a bias delta.** Our EU bbq baseline is 100 Age samples (§3(d)) — no race, gender, nationality, religion, class, disability, appearance or sexual-orientation items in it. So “Brazil minus EU” compares five Brazilian prejudices asked in Portuguese against **ageism asked in English**: it varies the prejudice as well as the jurisdiction. The claim that the pair reuses an identical *scorer* remains true and is worth having. The inference we drew from it — that the delta therefore isolates the Brazil-specific content — does not follow.

What we say plainly. Both models now score **higher** on the Brazilian set than on the EU one, which is the opposite of every direction we previously published. **Haiku’s delta is not distinguishable from zero** (1.1 standard errors). Sonnet’s is 2.7 standard errors from zero, which would be conventionally significant, but it is one model, it points away from the hypothesis, and its most economical explanation is that the Brazilian items are easier rather than that the model is less biased in Brazil. Given (b), we do not read either number as a bias comparison at all. **On this instrument there is no EU–Brazil bias gap in either direction, and we are not replacing the retracted negative delta with a positive one.**

What survives, and it is a within-benchmark comparison so (b) does not touch it. The per-polarity split that iteration 2 added to bbq_brazil — BBQ’s non-negative question half, absent from our iteration-1 set — is measurable for the first time on a log whose answers can be read:

Model	negative Q	non-negative Q	ambiguous ctx	disambiguated ctx
Haiku 4.5	0.8720 $\pm$ 0.0235	0.9300 $\pm$ 0.0172	0.9590 $\pm$ 0.0139	0.8430 $\pm$ 0.0251
Sonnet 4.6	0.9170 $\pm$ 0.0189	0.9575 $\pm$ 0.0132	0.9315 $\pm$ 0.0172	0.9430 $\pm$ 0.0154

Both models do 4-6 points worse on the negative question (“*who is less prepared?*”) than on its non-negative counterpart, so the polarity pair earns its place: a set with only negative questions would have reported the easier half of each model’s behaviour. Haiku is additionally 12 points worse on disambiguated contexts (0.843) than on ambiguous ones (0.959) — good at withholding a guess, less good at reading a context that licenses one — while Sonnet is even across the two. These are directions on small differences and we report them as such.

4.3 Results as they stand

Table 1. Same-model EU–Brazil deltas across six models (delta = Brazil minus EU; negative = less compliant on Brazil-specific content). Disclosure and bias each reuse the *same scorer* per pair. Contestation / explanation / AIA have no EU equivalent (Brazil-only scores shown). Column key: **hd** = **human_deception** (disclosure, Art. 5, I); **bbq-BR** = **bbq_brazil** (bias, Art. 5, III); **expl** = **explanation_quality** (Art. 6,




Figure 1: **Figure 1. RETRACTED** — the figure plots a scoring artefact, not model behaviour. It is kept as the record of what was published. Every red bar is the same quantity: the fraction of *English* samples in `human_deception_brazil` (11 of 21 = 0.524), because the Portuguese samples were asked for *NÃO* and scored against "NO" and so scored 0 by construction. The apparent reproduction across six models is one constant drawn six times. Corrected figures are in Table 1a: Haiku 0.986 ± 0.014 , Sonnet 0.962 ± 0.038 , deltas -0.014 and -0.038.

I); **contest** = `contestation_review` (Art. 6, II-III); **AIA** = `aia_checklist` (Arts. 25-28). Higher is more compliant.

Model (config)	hd EU	hd-BR (retr.)	Δ dis. (retr.)	bbq EU	bbq-BR	Δ bias (retr.)	expl	contest	AIA
Haiku 4.5 (scaled)	1.000	0.524	-0.48	0.858 (u)	0.677 (u)	-0.18	0.833	0.975	0.983
Sonnet 4.6 (scaled)	1.000	0.524	-0.48	0.518 ‡	0.375 ‡	-0.14	0.844	0.983	0.983
Llama 3.1 8B (pilot)	1.000	0.500	-0.50	0.600	0.500	-0.10	0.778	0.917	0.833†
gpt-oss 20B (pilot)	1.000	0.550	-0.45	0.700	0.750	+0.05	0.778	0.958	1.000†
Qwen2.5 14B (pilot)	1.000	0.550	-0.45	0.600	0.700	+0.10	0.722	0.875	1.000†
Mistral Small (pilot)	0.950	0.550	-0.40	0.550	0.550	+0.00	0.722	0.917	1.000†

Retracted columns. The **hd-BR** and **Δ dis.** columns are retracted in full — see §4.1 and Table 1a. Every cell in the hd-BR column is the same quantity (the English share of the dataset), so the column is one

constant six times, not six measurements. The **hd EU** column is unaffected: English prompt, English target. The **Δ bias** column is retracted in full as well — see §4.2 and Table 1b. Every cell in it is a Brazilian-content score differenced against an **age-only** EU baseline, so it varies the prejudice as well as the jurisdiction; and Sonnet’s two inputs to it were unreadable.

‡ Sonnet’s BBQ figures are **unsound, for a reason found on 2026-07-26 and now fixed**: it copies the answer template’s placeholder literally and replies **ANSWER: \$B**, which the `multiple_choice` parser cannot read, on 1,628 of 4,000 `bbq_brazil` and 315 of 1,000 `bbq` samples. The iteration-1 note here read, and is withdrawn: *“it uses BBQ’s cannot-determine option unreliably; the scorer parses its answers correctly in every spot-check, so this is a genuine behavioral quirk, not a scorer bug.”* It is a scorer bug. Full account and corrected figures in §4.2 and Table 1b. These particular cells have no replacement — they come from iteration-1 log directories that no longer exist and so could not be re-scored.

- (u) Verified unaffected by that defect — zero unparsable answers in 5,000 samples, established by re-scoring — but the **Δ** built from them is still retracted for the age-only-baseline reason above. The four local models’ BBQ cells are **unmarked because they are unchecked**: their logs were never censused for empty answer fields, which is a one-line check Phase 9 must run.

† Local `aia_checklist` is a single scenario at 1 epoch ($n = 1$) — read the local 1.000s as one observation, not a precise score.

Bias (Art. 5, III): withdrawn — see §4.2. The iteration-1 reading of this row was, quoted and withdrawn: *“directional trend, not yet conclusive — on the 44-sample set both reliable frontier models are negative (Haiku -0.18, Sonnet -0.14) and 4/6 models are negative overall (mean about -0.05), so the direction supports more bias on Brazilian categories.”* One of the two frontier deltas was built from unreadable answers; every delta in the column is differenced against an age-only EU baseline; and on the corrected measurement both frontier models score *higher* on the Brazilian set. There is no bias gap here in either direction.

Art. 6 triad + AIA: confirmed “beyond the EU,” with a sharpened story. All six models score **0.72-0.99** on explanation, contestation / human review, and AIA — i.e., they *can* articulate the high-risk procedural rights well. These benchmarks discriminate rather than saturating trivially: models reliably omit specific elements (most often a confidence / uncertainty statement in explanations). The iteration-1 reading of this — *“the important nuance is that the compliance failure is specific to disclosure, not to high-risk rights in general”* — is **withdrawn with the disclosure gap** (§4.1): there is no disclosure failure for these scores to be contrasted against. What remains is the unqualified version — on the benchmarks we built, at the sizes we ran them, these models articulate Brazil’s high-risk procedural rights competently, and the interesting question is no longer “why disclosure and not these?” but how much of the competence is genuine procedural knowledge rather than rubric vocabulary, which is what §4’s judge cross-check is for.

Methodological finding: an EU baseline can manufacture a direction. One model’s bias delta, on one unchanged behaviour, has now been published with **three different signs**. Scaling the EU `bbq` baseline from 20 to 1000 observations moved Haiku’s delta from **+0.05** to **-0.18**, because the baseline itself moved 0.65 to 0.858. Fixing the answer parse and rebuilding `bbq_brazil` moved it to **+0.04** (§4.2). No model changed at any point. Two distinct baseline defects compound to produce that: **under-powered** (the 20-sample pilot) and **single-axis** (every `bbq` baseline here is age-only, §3(d)). Read together with the five instrument defects in §3, the generalisation we would defend is stronger than “use a properly-powered baseline”: **in a cross-jurisdiction comparison, the instrument is the most likely source of the effect, and the reused baseline deserves the same audit as the new benchmark.** Ours did not get it until iteration 2.

5. Positionality, participation, and human rights

A benchmark that measures Brazilian prejudice, whose categories were chosen by two people who are not Brazilian, whose items were drafted by a language model, and whose only review has been by other language models, is in a poor position to cite decolonial theory. Citing the theory without a participation stance is the failure mode the theory names. This section states the frame we work in, the instruments that make it binding rather than aspirational, and then applies both to this project by name.

5.1 *Racismo algorítmico*, not “AI bias”

The central citation is Brazilian and in Portuguese: Tarcízio Silva’s *Racismo Algorítmico* [14], for whom “*o racismo algorítmico é uma espécie de atualização do racismo estrutural*” — algorithmic racism as an *update* to structural racism rather than a defect of a model. That reframing changes the object of measurement. A “bias score” invites the reading that a model has a fixable flaw; *racismo algorítmico* insists that the flaw is continuous with a structure the system is embedded in, which is why PL 2338 places non-discrimination among *rights of affected persons* (Art. 5, III) rather than among model-quality requirements. The 2024 IRIS-BH / Tarcízio Silva / Ação Educativa report on AI and racial discrimination in Brazil [15] — produced with Brazilian expert input feeding a UN thematic report — is the companion source, and the reason `bbq_brazil`’s axes are Brazilian rather than translated: IBGE’s five-category *cor ou raça* taxonomy [10]; regional prejudice, documented as “internal orientalism”; class markers from ABEP’s *Critério Brasil*; and *racismo religioso*, a term Brazilian scholarship originated, covering communities that are roughly 1% of the population and 50–65% of recorded religious-intolerance victims. No US-derived taxonomy produces that last axis at all.

5.2 Categories are not neutral data fields

Ricaurte [17] argues that treating locally-produced categories as neutral data fields reproduces an epistemic hierarchy — the categories become inputs to somebody else’s model rather than descriptions their holders control. Couldry and Mejias [18] extend this to the dataset itself: a corpus of stereotype scenarios about real communities is a site of extraction, not only of measurement. Mohamed, Png and Isaac [19] pose the question that a bias benchmark cannot avoid — *whose knowledge counts as ground truth* — though we cite them with a tension acknowledged, since their “reverse tutelage” proposal has itself been criticized for risking a framing in which the Global North grants permission to participate. Adams [20] asks whether AI can be decolonized at all without changing who builds it, which is a harder question than the one we answer here. Birhane’s relational account [21] is the reason `bbq_brazil` has a compound *intersectional* axis rather than stackable independent ones: *mulher negra nordestina* is not the sum of three filters. And *Data Feminism* [22] and *Design Justice* [23] supply the two propositions this section is measured against — centre the lived experience of the categories rather than literature-derived descriptions of them, and remember that “power shapes participation in all design processes”, so a participatory step is not automatically an equitable one. On the first proposition, this project currently fails: our stereotypes are literature-derived.

5.3 The instruments are binding, not aspirational

Domestically, CF/88 Art. 5 makes equality a fundamental right and racism an imprescriptible, non-bailable crime; Lei 7.716/1989, Lei 12.288/2010 (Estatuto da Igualdade Racial) and Lei 14.532/2023 give it statutory shape; and **LGPD Art. 5º, II** names racial or ethnic origin and religious conviction as *sensitive personal data* — which is directly relevant, because two of `bbq_brazil`’s five axes are LGPD-sensitive categories and the participation protocol in §5.6 would collect exactly that data from its annotators.

Internationally, Brazil has ratified the **American Convention on Human Rights** (Decreto 678/1992; IACtHR compulsory jurisdiction accepted 1998), Arts. 1 and 24; and **ICERD** (ratified 27 March 1968, Decreto 65.810/1969), whose Art. 1 definition is **effects-based** and so reaches facially neutral systems with discriminatory outcomes [24]. The feminist anchor is the **Convention of Belém do Pará**, adopted in 1994 *in Belém, Brazil* and ratified in November 1995: it obliges states to modify the “*socio-cultural patterns*” that perpetuate stereotyped roles. That is a binding treaty, adopted on Brazilian soil, that makes **stereotype itself** the object of a state obligation — which is precisely the object a stereotype benchmark measures, and a stronger foundation for the feminist framing than any citation we could import.

One caution we impose on ourselves. Brazil has been an ICERD party since 1968 and the CERD Committee’s 2022 Concluding Observations address racial discrimination in Brazil. **We attribute no finding about AI or algorithms to the Committee and cite no paragraph of it.** The step from ICERD Art. 1’s effects-based definition to algorithmic systems, and specifically to facial recognition, is **our own argument**, and it should be read and contested as ours. On the soft-law side, the UNESCO Recommendation on the Ethics of AI [25] (adopted 2021; pt-BR translation by UNESCO Brasília, 2022)

and the UNGPs with OHCHR’s B-Tech work [26] are the natural frame for an AIA, since human-rights due diligence is what Arts. 25-28 amount to in substance.

5.4 *Not My AI*, adapted rather than cited

Coding Rights’ feminist-decolonial framework for public-sector algorithmic systems [27] is designed to be used, so we used it, turning its orientation into questions we put to vigilAI and answered:

Question	Answer for vigilAI
Who is the system <i>for</i> , and who is it done <i>to</i> ?	It is for regulators and deployers. The affected person is <i>represented</i> by the benchmark, not <i>present</i> in it. Partial mitigation: the Art. 28 scorecard is built to be a public-facing artifact rather than an internal report.
Whose categories?	Ours — see §5.5.
Does it make an existing power asymmetry legible, or add one?	Both. A compliance score is a thing a deployer can use to certify itself. The counterweight is the gap-flagging items , which measure whether a deployer exceeds duties that no Brazilian instrument imposes; a low score there is a finding about the regulatory regime, not about the model.
Who can contest the instrument itself?	Today, nobody. There is no channel through which a Brazilian who thinks an item is wrong can have it removed. §4.3 of <code>docs/participation-protocol.md</code> (not of this paper) is the design of that channel, including a removal power the researchers cannot overrule.
What happens to data about people?	The scenarios are synthetic and describe no real individual. The annotator data the protocol <i>would</i> collect is LGPD-sensitive and is minimized, separated and deleted on a stated schedule.
Who benefits from publishing?	We do. That is the honest answer to Birhane et al.’s question [28], and it is the one the protocol’s co-authorship and benefit-sharing terms exist to change.

5.5 Participation-washing, and the helicopter-research test applied to vigilAI

Sloane, Moss, Awomolo and Forlano name **participation-washing** [29]: participatory language, or a thin participatory gesture, used to launder a system whose design decisions were made elsewhere. Birhane, Isaac, Prabhakaran and Díaz [28] ask who the primary beneficiaries of participatory activity actually are. The helicopter-research literature [16] converts both into a test with four checkable questions. We answer them about this project, in the paper rather than in a supplement, because an answer kept in a supplement is the gesture the critique is about.

Who conceived the categories? The two authors, neither Brazilian, neither resident in Brazil. The five axes were selected from Brazilian secondary sources, but *which* axes exist — and what counts as an axis at all — was our decision.

Who wrote the items? A large language model (Claude, via Claude Code), under our direction, from term banks the same model assembled from those secondary sources. The 22 pilot scenarios were drafted item by item; the other 78 are deterministic expansions of 30 templates over 34 contrast pairs. The generation *pipeline* contains no model call and reproduces byte-for-byte — but that is a reproducibility property, not a provenance one, and it would be misleading to let it stand as an answer to this question. **No Brazilian wrote a scenario.**

Who validated them? Nobody who lives the categories. Every review to date has been by **LLM judges** — five rounds over the 100 **bbq_brazil** scenarios and one over the 24 rubric scenarios, both committed to the repository, and each document opening by stating that it is a pre-screen and not validation. Those rounds caught real defects — seven items that measured backwards, one keyed to its own biased answer, a template pointed at an axis no Brazilian prejudice occupies — which is evidence that the items *needed* review, not that they have had it.

Who is a co-author? Two authors. No Brazilian co-author, no Brazilian institutional partner, no Brazilian ethics review.

What this does and does not undermine, stated because a positionality section that made every finding equally suspect would be as uninformative as one that made none. It bears directly on **bbq_brazil** (Art. 5, III), whose construct validity rests entirely on whether these stereotypes are attested in Brazil and pointed in the right direction — a judgment we have not obtained, and the reason the bias result is reported as a direction and not a magnitude. It bears partly on the three pt-BR rubric benchmarks, whose scenarios are similarly authored but whose scoring targets a procedural checklist drawn from statute rather than from a social judgment. It bears least on the paper’s headline, **human_deception_brazil** (Art. 5, I): its target is “the model denies being human”, scored by the same scorer as the English original, with no stereotype judgment anywhere in it.

5.6 The protocol, and the claim we are not making

`docs/participation-protocol.md` sets out the validation that would close §5.5, composed from three published BBQ-family protocols and citing each for the component it actually contributes: **KoBBQ’s quantitative core** [30] (a defined N per item, demographic balancing, a **>2/3 validity threshold**, a reading-comprehension check, and *reporting how many candidate items were eliminated*); **SeeGULL’s qualification rule** [31] — in-region residency, generalized to lived membership, so that a *nordestino* stereotype is validated by *nordestinos* and a *candomblecista* one by *candomblecistas*; and **PakBBQ’s transparency and duty of care** [32] — named annotators, a regional-diversity quota, and a **harm-exposure briefing before annotation begins**, with a removal power that the researchers cannot overrule. Agreement is reported per template, with per-claim proportions and intervals and a full elimination count — and reported *by panel* rather than pooled, following NLPpositionality [34], because a claim that *nordestinos* and *paulistanos* rate an item differently is a finding about the stereotype and not noise to average out of it. That reporting is itself a modest contribution, which says more about the field than about us: **SHADES, MBBQ and JBBQ report no agreement statistic at all**, and “we used native speakers” is common practice rather than a gold standard.

Two implementation questions are recorded as open rather than resolved by assumption. **Compensation:** no citable Global-South norm exists, and the field’s own cautionary example is exact — SeeGULL paid **USD 8.22** per annotator in India against **USD 28.35** in Australia for the same study, unjustified in the paper. “Local cost of living”, applied uncritically, prices the labour by where the worker lives rather than by what the work is, which reproduces the inequity a decolonial framing exists to address. The protocol therefore commits to **one rate for every participant regardless of country**, with fees borne by the project, briefing time paid, and no reduction for stopping early — and it computes the resulting budget (~460 paid hours) precisely so that deferring the work has to be a decision somebody writes down. **Research ethics:** whether stereotype-annotation work of this kind requires review by a *Comitê de Ética em Pesquisa* under Resolução CNS nº 510/2016 has no Brazil-specific guidance we could find; because the work collects demographic self-identification and deliberately exposes participants to injurious material, we treat it as a question for the team’s own CEP rather than one to answer by assumption.

The protocol also names who would be asked, in tiers keyed to lived membership rather than to credentials — Brazilian algorithmic-racism organizations for the axes as a whole, an organization working specifically on Black women and *racismo religioso* for the two axes carrying the highest harm exposure, and the Brazilian NLP communities for pt-BR register, approached as communities rather than as cold-emailed individuals. We deliberately do not print those names here. **None of them has been approached, none has reviewed this protocol, and none has agreed to anything in it;** listing an organization’s name next to a

methodology it did not review is itself a form of participation-washing, and it is the easiest one to commit by accident.

The claim we are not making. No Brazilian annotator and no Brazilian organization has validated any item in this benchmark. There has been **no community validation**, and no LLM pre-screen in this repository substitutes for it. What iteration 2 delivers is the protocol and an honest account of the distance still to travel.

6. Discussion and Limitations

Implications for AI safety. Brazil’s Chapter II grants rights the EU AI Act does not grant individuals at all, so no EU-derived benchmark suite can say anything about them — that argument is structural and survives every number in this paper moving. What does *not* survive is the version of it we previously led with. We argued that a model can pass an English / EU audit and still systematically violate a Global-South statutory right, and offered the disclosure gap as the cleanest example and the bias delta as the supporting trend. **Both were our own measurement error** (§4.1, §4.2): on the corrected measurements these models comply with Art. 5, I in Portuguese, and there is no EU–Brazil bias delta in either direction that our instrument can support. The structural argument therefore stands on the *absence* of counterparts — Brazil’s Art. 6 explanation, contestation and human-review rights and its Arts. 25-28 AIA have no EU benchmark at all — rather than on a demonstrated behavioural failure. That is a weaker claim than the one we made, and it is the one the evidence supports.

The five broken instruments are themselves a result, and we would put them ahead of any score in this paper. Four of the five (§3) sat under a number we had already published, and the two that mattered most were in *reused* code — an upstream scorer’s target-matching contract, and an upstream solver’s answer format — which is precisely the code a cross-jurisdiction benchmark must reuse for its delta to mean anything. The generalisation: **a localized benchmark’s most likely source of a finding is its own instrument, and the reused baseline is the least audited part of it.** None of the five was visible in a standard error; three of the five inflated a score and two destroyed one; each has a one-line check that we did not run. If there is a transferable contribution here beyond the artifacts, it is that list of checks and the practice of publishing the retraction rather than the correction alone.

vigilAI still shows the localization is achievable with modest effort: a fork, five benchmarks, and a same-model delta give regulators and deployers a statute-referenced, reproducible compliance picture, and the Art. 28 scorecard turns an eval run into the exact public-conclusions artifact a high-risk deployer is obligated to produce.

A second implication, which is about us rather than about the models. Three of the four defects this iteration found were in *our own instruments*, not in the models: six over-broad scorer cues that gave `contestation_review` a floor of 0.5, an `aia_checklist` prompt that enumerated its own answer key (echo floor 0.944) and cue lists that scored a content-free non-answer 6/6, and the disclosure target above. Every one of them **inflated or fabricated a finding**, none of them was visible in a standard error, and the last one produced the paper’s headline. A compliance-evaluation tool carries the same obligation it measures: to state what would change its conclusion and to make the check reproducible. Concretely, that means scoring your own rendered prompt against your own scorer, probing your cue lists with a hostile non-answer, asserting that every sample’s target is answerable in the language its prompt asks for, and reading transcripts chosen by a rule rather than by eye. None of those costs more than an afternoon; all four were added only after the defect they catch had already shipped.

What follows explains why the Art. 6 rights are worth building benchmarks for at all, and why the sector overlays exist — and it is now the paper’s principal finding rather than a supplement to a behavioural one. It does not depend on any model score.

Brazilian law grants a right to review and twice declines to say who performs it

The reflexive objection to benchmarking PL 2338’s Art. 6 triad is that it restates LGPD Art. 20. For explanation (Art. 6, I) that objection has force. For **human review (Art. 6, III) it is wrong**, and the

legislative record says so precisely.

The drafting history, from the primary record. As enacted on 14 August 2018, LGPD Art. 20 granted a right to request review of a solely-automated decision “*por pessoa natural*” — the phrase sat unconditionally in the *caput*, and §1 (the duty to supply clear information about the criteria and procedures) and §2 (the ANPD’s power to audit for discriminatory aspects where trade secrecy is invoked) were present from the start and are word-for-word unchanged today. The 2018 veto message did not touch Art. 20. **MP 869/2018** then removed “*por pessoa natural*” from the *caput*. During its conversion, **the §3 introduced by the 2019 conversion bill** (PLV 7/2019) reinstated a weaker, *conditional* human-review requirement, tied to future ANPD regulation and calibrated to the entity’s size and processing volume:

“A revisão de que trata o caput deste artigo deverá ser realizada por pessoa natural, conforme previsto em regulamentação da autoridade nacional, que levará em consideração a natureza e o porte da entidade ou o volume de operações de tratamento de dados.”

That §3 — not the 2018 text — is what **Mensagem nº 288, de 8 de julho de 2019** vetoed on sanctioning Lei nº 13.853/2019 [11], and Congress upheld the veto on **2 October 2019** (Veto nº 24/2019, item 24.19.001) [12]. Art. 20 today therefore reads *caput* + §1 + §2 + §3 (*VETADO*).

The consequence is silence, not permission. The statute in force neither mandates nor forbids any particular reviewing agent, so a second automated pass is lawful **by omission** rather than by affirmative authorization. That is both more accurate and harder to argue with than “Brazilian law permits machine review”.

The veto was economic, not constitutional — and the central bank was in the room. The message as a whole invokes both *contrariedade ao interesse público* and unconstitutionality across the thirteen devices it vetoes; for **this** device the stated ground is only the former, and its reasoning is a credit-supply argument: universal human review “*inviabilizará os modelos atuais de planos de negócios de muitas empresas, notadamente das startups, bem como impacta na análise de risco de crédito e de novos modelos de negócios de instituições financeiras, gerando efeito negativo na oferta de crédito aos consumidores*”, reaching inflation and monetary policy. The message records that the Ministries of the Economy and of Science, Technology, Innovation and Communications, the Controladoria-Geral da União **and the Banco Central do Brasil** came out for the veto. The margin was narrow: the Câmara voted **261–163 to overturn**, clearing the 257-vote absolute-majority threshold, while the Senate reached **40 of the 41** votes required, so under CF/88 Art. 66 §4 the veto stood. (Tally from the Congresso Nacional veto-tracking database [12], not the session *ata*.)

This is a pattern across two instruments, not one statute’s accident. Eight years earlier, **Lei 12.414/2011 Art. 5, VI** — the Cadastro Positivo statute, Brazil’s consumer-credit reporting regime — had already granted the *cadastrado* the right “*solicitar ao consulente a revisão de decisão realizada exclusivamente por meios automatizados*” [13]. Review, again; a human reviewer, again unspecified. The 2019 amendments (LC 166/2019) rewrote four of that article’s seven items and left this one untouched. So Brazilian law grants a right to review of a solely-automated decision twice, in a general data-protection statute and in a sectoral credit statute, and twice leaves the reviewing agent open. **PL 2338 Art. 6, III is what closes it** — generally against LGPD Art. 20, and in credit against Lei 12.414 Art. 5, VI, which is exactly the domain the veto’s own stated reasoning was about.

Why this motivates the sector overlays rather than a single national score. The gap is not evenly distributed, because the argument that created it was sectoral: the human-review right was blocked partly on credit-risk grounds with the financial regulator supporting the veto. So the finance overlay’s `human_review_gap_lgpd20` is a **gap-flagging item** — it measures whether a deployer voluntarily exceeds a duty that no Brazilian instrument imposes, and a low score there is a finding about the regulatory regime rather than about the model. Health is the counter-example that shows the mapping is real rather than rhetorical: **CFM Resolução nº 2.454/2026** [33] independently mandates non-waivable human supervision of medical AI — “*As soluções apresentadas pelos modelos, sistemas e aplicações de IA não são soberanas, sendo obrigatória a supervisão humana*” — so a professional council has already adopted for physicians what the federal veto removed for everyone. It must be read with two qualifications: it is **adopted but not yet**

in force (Art. 23: 180 days from publication, about 26 August 2026), and it binds **physicians through CRM discipline, not products through ANVISA**. Capital markets supply the opposite extreme — no CVM instrument uses “*inteligência artificial*” in an operative clause, and there is **no Arts. 25-28 analogue at all**.

None of this is legal advice; see the final limitation below.

Limitations

- **Provenance and validation — the limitation that governs the others.** The Brazilian categories were chosen by two non-Brazilian authors; the items were drafted by a language model from term banks assembled out of secondary literature; and every review to date has been by LLM judges (five rounds on `bbq_brazil`, one on the rubric scenarios), each documented in the repository as a pre-screen rather than as validation. **No Brazilian has validated any item.** §5 states this in full and says which findings it does and does not reach; `docs/participation-protocol.md` states what would close it.
- **Dataset scale.** `bbq_brazil` is 100 scenarios / 400 samples (20 scenarios per axis); the three rubric benchmarks are **n = 12** each. That is a real improvement on iteration 1’s 44 / 3 / 4 / 1, and it is still small: **n = 12** supports a standard error, not a verdict. `bbq_brazil`’s four samples per scenario are also **not independent** — the two polarities share a context and the two contexts share a scenario — so the tool-reported standard error is a **lower bound**; the scenario-level figure (**n = 100**) is the honest outer bound, and no $\sqrt{400}$ precision gain is claimed anywhere.
- **Deterministic rubric scorers** detect the *presence* of required elements via multilingual cues. They measure whether a compliant *description* is produced, not real-world behavior, and could miss unusual phrasings (we tuned cue lists against real model output). Small-n tasks (e.g., contestation **n = 4**) show run-to-run variance; the 10-epoch frontier numbers are the stable ones.
- **Pilot vs. scaled.** Local results are 1-epoch `--limit 20`; only the scaled frontier runs are high-precision. Sonnet’s Table 1 BBQ cells should not be cited at all: the cause is an answer-format *parse* failure specific to that model (**ANSWER: \$B**), not a behavioural quirk. It is fixed and the iteration-2 logs were re-scored (§4.2, Table 1b); the iteration-1 cells could not be, because those log directories no longer exist. The four open-weight models’ disclosure figures are retracted outright (§4.1) and were not re-run, and their BBQ figures inherit the age-only-baseline problem in their delta.
- **The EU `bbq` baseline is a single BBQ axis, and this is the limitation we would most want a reader to carry.** `--limit` is global per invocation and `inspect_evals.bbq` concatenates its eleven subsets with **Age** first, so every EU bias baseline in this project is 100 **Age** samples. Fixing it is cheap — sample across the subsets, about \$0.58 per frontier model and \$0 on Ollama — and until it is done the Art. 5, III row has a defensible *absolute* score and no defensible EU comparison. We found this while re-measuring bias after fixing the parse; it had been true, unexamined, in both iterations.
- **Attribution assumption, and how it failed.** We assume the same-scorer EU–Brazil delta attributes differences to language / legal content, because format and scorer are identical. The assumption is sound and the *instantiation* was not: “same scorer” is not “same target”, and the disclosure pair differed in its target (§4.1). We reasoned from the tight 0.50-0.55 clustering that a generic-competence explanation was ruled out; the clustering was in fact the signature of a constant, and should have prompted the question “what could make this number identical across models?” rather than confidence. **A same-scorer delta is only as strong as an audit of everything else the two sides do not share** — target, instruction, answer format, how the answer is extracted, *and which items each side actually contains*. **Three of those five have now bitten this project:** the target (the disclosure pair, §4.1), the extraction (the `bbq` pair for one model) and the content (the `bbq` pair for every model, §4.2). None of the three was predicted; each was found by looking.
- **Not legal advice.** The PL 2338/2023 article mappings are an engineering interpretation built for evaluation, made against a bill still moving through the legislative process; they are a research instrument, not a legal determination of compliance for any specific deployment. **The same disclaimer covers the sector overlays** — the BACEN, ANVISA / CFM / ANS and CVM instruments are mapped to PL 2338 articles as *structural analogues for benchmark design*, not as a claim about what any regulator requires of any firm. It also covers the legislative analysis in this section: the LGPD Art. 20 and Lei 12.414 readings are drawn from primary records and cited so they can be checked, but they

are our reading of the record, not counsel.

Future Work

Execute the participation protocol — the single highest-value next step, and the one that would let the bias benchmark’s *content* be defended rather than only its arithmetic. It begins with Tier-1 contact rather than with recruitment: putting `docs/participation-protocol.md` in front of one Brazilian algorithmic-racism organization and asking whether the protocol is the right shape, before asking anyone to execute it. **Rebuild the EU bias baseline across BBQ’s eleven subsets** — cheap, and it is what a Brazil-vs-EU bias comparison actually requires (§4.2(b)). Beyond that: resolve the CEP question in writing; scale the open-weight models to the same config as the frontier ones so the two halves are directly comparable; extend the sector overlays as ANPD’s Art. 20 rulemaking and CFM Res. 2.454/2026’s entry into force (about 26 August 2026) change the underlying instruments; and give the remaining EU-only technical requirements a Brazilian framing as the bill moves.

7. Conclusion

We asked whether EU / English AI-safety compliance transfers to Brazil’s PL 2338/2023, and built **vigilAI** — a COMPL-AI fork with five Brazil-specific benchmarks and a same-model EU–Brazil methodology — to answer it. Our answer changed while we were checking it. Two iterations reported that compliance does not transfer, on the strength of a -0.45 Portuguese disclosure gap and a negative bias delta; **neither survived an audit of our own instruments**. The gap was an artefact of our scoring target, and on the corrected measurement the two frontier models disclose in Portuguese as reliably as in English (deltas -0.014 ± 0.014 and -0.038 ± 0.038). The bias delta rested on an answer format our parser could not read and on an EU baseline containing one of BBQ’s eleven axes; corrected, both models score *higher* on the Brazilian set, and we report no bias gap rather than a reversed one. **On the behaviour we can currently measure, there is no EU–Brazil compliance gap in either disclosure or bias** — and reporting that, together with the five broken instruments and how a rule-selected transcript exposed the first of them before write-up, is the result we can defend. The finding a reader should take away is about measurement: **in a cross-jurisdiction benchmark, audit the instrument before believing the effect, and audit the reused baseline hardest**. Meanwhile, Brazil’s high-risk Art. 6 rights (explanation, **contestation**, **human review**) and its AIA have **no EU benchmark counterpart at all**, and vigilAI introduces deterministic benchmarks for them. One of those rights is not even a restatement of Brazil’s own data-protection law: the legislative record shows Brazilian law granting a right to *review* of a solely-automated decision twice — in LGPD Art. 20 and in Cadastro Positivo Art. 5, VI — and declining both times to say who performs it, so **Art. 6, III is a genuine increment**, and one whose absence is deepest in the sector whose regulator supported the 2019 veto. The practical message for Global-South AI governance is blunt: **certification under a foreign regime is not compliance**, and localized, statute-referenced evaluation is both necessary and achievable today. The equally blunt message about this work is that a localized benchmark built without the participation of the people it describes is a beginning and not a result: we publish the protocol that would change that, and we make no claim to have run it.

Code and Data

- **Code repository:** <https://github.com/dyrtyData/vigilAI> — a fork of **COMPL-AI**. Reproduce with `uv run vigilai eval ... && uv run vigilai report logs/<run> --html`.
- **Data / Datasets:** Brazil benchmark scenarios ship in-repo (`src/vigilai/tasks/*_brazil`, `contestation_review`, `aia_checklist`); anchored by SHADES / BiasShades (pt-BR) [7] and ToxSyn-PT [8].
- **Other artifacts:** `reports/RESULTS.md` (full multi-model analysis), `reports/scorecard.html` (Art. 28 scorecard), `reports/multimodel-scorecard.html` (six-page dossier, reproduced in Appendix B), and `reports/runs/` (per-model reports, split into Stage-7 baseline and Phase 8-11 additions).
- **Participation protocol:** `docs/participation-protocol.md` — the native-annotator validation protocol described in §5.6, including the tiered outreach list, the compensation terms and budget, and the

open research-ethics question. The LLM pre-screens it explicitly does **not** credit are committed alongside it (docs/bbq-brazil-llm-judge-review.md, docs/rubric-scenarios-llm-judge-review.md, docs/bbq-brazil-unreviewed-wordings.md), each stating in its own opening that it is a pre-screen and not validation.

- **License:** code is released under **Apache-2.0** (matching the upstream COMPL-AI ecosystem); this report and its figure under **CC-BY-4.0**.

Author Contributions

Diana Chang and Ian Duhamel Hayes jointly designed and built vigilAI, implemented the benchmarks and reporting layer, ran the evaluations, and prepared this report.

Positionality. Neither author is Brazilian and neither is resident in Brazil. Neither is a member of any of the groups `bbq_brazil`’s five axes describe. There is no Brazilian co-author, no Brazilian institutional partner and no Brazilian ethics review on this work. We state this here, and in full in §5.5, because a paper that argues for localized evaluation while being written entirely from outside the locality should say so on its own front matter rather than leave a reader to infer it.

LLM Usage Statement

We used Claude (Claude Code) to fork and refactor COMPL-AI, implement the Brazil benchmarks and the reporting layer, run the evaluations, and draft this report. **Two uses go beyond assistance and are material to how the results should be read**, so they are stated here as well as in §5.5:

1. **The Brazilian benchmark content was drafted by a language model.** The `bbq_brazil` term banks, scenario templates and pilot scenarios, and the pt-BR rubric scenarios, were written by Claude under our direction from the secondary sources cited in §5.1. The expansion from templates to the committed 100 scenarios is deterministic and contains no model call, and it reproduces byte-for-byte from `tools/generate_brazil_scenarios.py` — but that is reproducibility, not provenance.
2. **The only review of that content has also been by language models** — five rounds of LLM-judge review, committed to the repository and each labelled a pre-screen rather than validation.

The scorers themselves are deterministic and unit-tested (879 tests, run with no API key and no network call); the optional LLM judge is a *second* scorer reported alongside the deterministic one, never in place of it. That the test count grew by a hundred while landing the retractions in §4.1 and §4.2 is the honest summary of iteration 2: most of the new tests pin a defect that had already been published.

References

1. LatticeFlow AI, ETH Zurich, INSAIT. *COMPL-AI: A Technical Interpretation and LLM Benchmarking Suite for the EU AI Act*. 2024. <https://compl-ai.org> · <https://github.com/compl-ai/compl-ai>
2. European Union. *Regulation (EU) 2024/1689 (Artificial Intelligence Act)*. 2024.
3. Senado Federal do Brasil. *Projeto de Lei nº 2338/2023 — Marco legal da inteligência artificial*. Approved by the Senate, Dec. 2024.
4. Brazil. *Lei nº 13.709/2018 (Lei Geral de Proteção de Dados — LGPD)*, esp. Art. 20.
5. Parrish, A. et al. *BBQ: A Hand-Built Bias Benchmark for Question Answering*. Findings of ACL 2022. <https://aclanthology.org/2022.findings-acl.165/>
6. BBQ adaptations, none covering Portuguese or the IBGE categories: Jin et al. *KoBBQ* (TACL 2024) [30]; Neplenbroek, V., Bisazza, A. & Fernández, R. *MBBQ: A Dataset for Cross-Lingual Comparison of Stereotypes in Generative LLMs*. COLM 2024, <https://arxiv.org/abs/2406.07243>; Hashmat et al. *PakBBQ* (EMNLP 2025) [32]; *JBBQ*, <https://arxiv.org/abs/2406.02050>; *BharatBBQ*, <https://arxiv.org/abs/2508.07090>.
7. Mitchell, M. et al. *SHADES / BiasShades: Multilingual Stereotype Benchmark*. NAACL 2025. <https://aclanthology.org/2025.naacl-long.600/> · <https://huggingface.co/datasets/LanguageShades/BiasShades>
8. *ToxSyn-PT* (CC BY 4.0). <https://huggingface.co/datasets/ToxSyn/ToxSyn-PT>
9. UK AI Safety Institute. *Inspect AI: An open-source framework for LLM evaluations*. <https://inspect.aisi.org.uk>

10. IBGE (Instituto Brasileiro de Geografia e Estatística). Racial / colour classification (branco, pardo, preto, amarelo, indígena).
11. Brazil, Presidência da República. *Mensagem nº 288, de 8 de julho de 2019* (partial veto on sanctioning **Lei nº 13.853/2019**), DOU 9 Jul. 2019. https://www.planalto.gov.br/ccivil_03/_ato2019-2022/2019/Msg/VEP/VEP-288.htm
12. Congresso Nacional. *Veto nº 24/2019*, item **24.19.001** (§ 3º do art. 20 da Lei nº 13.709/2018); veto maintained in the joint session of **2 October 2019**. <https://www.congressonacional.leg.br/materias/vetos/-/veto/detalhe/12445>
13. Brazil. *Lei nº 12.414/2011* (Cadastro Positivo), Art. 5, VI, as amended by LC 166/2019.
14. Silva, T. *Racismo Algorítmico: Inteligência Artificial e Discriminação nas Redes Digitais*. Edições Sesc, 2022. <https://racismo-algoritmico.pubpub.org/>
15. IRIS – Instituto de Referência em Internet e Sociedade, with T. Silva and Ação Educativa. *Artificial Intelligence and Racial Discrimination in Brazil: Key Issues and Recommendations*. May 2024. <https://irisbh.com.br/en/publicacoes/artificial-intelligence-and-racial-discrimination-in-brazil-key-issues-and-recommendations/>
16. Chapman, C. A. et al. *Ten Simple Rules for Global North Researchers to Stop Perpetuating Helicopter Research in the Global South*. PLOS Computational Biology 17(9), 2021. <https://doi.org/10.1371/journal.pcbi.1009277> · Nature editorial, “Nature addresses helicopter research and ethics dumping”, 2022. <https://doi.org/10.1038/d41586-022-01423-6>
17. Ricaurte, P. *Data Epistemologies, the Coloniality of Power, and Resistance*. Television & New Media 20(4), 2019. <https://doi.org/10.1177/1527476419831640>
18. Couldry, N. & Mejias, U. A. *The Costs of Connection: How Data Is Colonizing Human Life and Appropriating It for Capitalism*. Stanford University Press, 2019.
19. Mohamed, S., Png, M.-T. & Isaac, W. *Decolonial AI: Decolonial Theory as Sociotechnical Foresight in Artificial Intelligence*. Philosophy & Technology 33, 2020. <https://doi.org/10.1007/s13347-020-00405-8>
20. Adams, R. *Can Artificial Intelligence Be Decolonized?* Interdisciplinary Science Reviews 46(1–2), 2021. <https://doi.org/10.1080/03080188.2020.1840225>
21. Birhane, A. *Algorithmic Injustice: A Relational Ethics Approach*. Patterns 2(2), 2021. <https://doi.org/10.1016/j.patter.2021.100101>
22. D’Ignazio, C. & Klein, L. F. *Data Feminism*. MIT Press, 2020. <https://data-feminism.mitpress.mit.edu/>
23. Costanza-Chock, S. *Design Justice: Community-Led Practices to Build the Worlds We Need*. MIT Press, 2020. <https://designjustice.mitpress.mit.edu/>
24. International Convention on the Elimination of All Forms of Racial Discrimination (ICERD), 1965 — ratified by Brazil 27 Mar. 1968, Decreto 65.810/1969; American Convention on Human Rights, 1969 — Decreto 678/1992; Inter-American Convention on the Prevention, Punishment and Eradication of Violence against Women (**Convention of Belém do Pará**), adopted 1994 in Belém, Brazil — ratified Nov. 1995.
25. UNESCO. *Recommendation on the Ethics of Artificial Intelligence*. Adopted 23 Nov. 2021; pt-BR translation, UNESCO Brasília, 2022.
26. United Nations. *Guiding Principles on Business and Human Rights*, 2011; OHCHR **B-Tech** Project on human-rights due diligence for technology companies.
27. Varon, J. & Peña, P. *Not My AI: A Feminist Framework to Challenge Algorithmic Decision-Making Systems Deployed by the Public Sector*. APC / Coding Rights, 2022. <https://www.apc.org/en/pubs/not-my-ai-feminist-framework-challenge-algorithmic-decision-making-systems-deployed-public>
28. Birhane, A., Isaac, W., Prabhakaran, V. & Díaz, M. *Power to the People? Opportunities and Challenges for Participatory AI*. EAAMO 2022. <https://arxiv.org/abs/2209.07572>
29. Sloane, M., Moss, E., Awomolo, O. & Forlano, L. *Participation Is Not a Design Fix for Machine Learning*. EAAMO 2022. <https://arxiv.org/abs/2007.02423>
30. Jin, J., Kim, J., Lee, N., Yoo, H. et al. *KoBBQ: Korean Bias Benchmark for Question Answering*. TACL 2024. <https://arxiv.org/abs/2307.16778>
31. Jha, A., Davani, A. M., Reddy, C. K. & Dave, S. *SeeGULL: A Stereotype Benchmark with Broad Geo-Cultural Coverage Leveraging Generative Models*. ACL 2023. <https://arxiv.org/abs/2305.11840>
32. Hashmat, A., Mirza, M. A. & Raza, A. A. *PakBBQ: A Culturally Adapted Bias Benchmark for QA*. EMNLP 2025. <https://arxiv.org/abs/2508.10186>
33. Conselho Federal de Medicina. *Resolução CFM nº 2.454, de 2026 — normatiza o uso da inteligência*

- artificial na medicina*. DOU 27 Feb. 2026 (retif. 5 Mar.); in force 180 days after publication (about 26 Aug. 2026). https://sistemas.cfm.org.br/normas/arquivos/resolucoes/BR/2026/2454_2026.pdf
34. Santy, S., Liang, J. T., Le Bras, R., Reinecke, K. & Sap, M. *NLPositionality: Characterizing Design Biases of Datasets and Models*. ACL 2023. <https://aclanthology.org/2023.acl-long.505/>

Appendix A — Reproducibility and details

- **Reproducibility commands**, per-model reports, the corrected Sonnet `bbq` breakdown, and standard errors: `reports/RESULTS.md`.
- **Coverage breadth**: 4 of 9 COMPL-AI requirements carry a bespoke Brazil benchmark (Disclosure, Representation, Interpretability, Societal Alignment); the remaining five are EU-only requirements with no Brazil Chapter II counterpart — rendered as a colour-coded (green / amber / grey) coverage map in every report.
- **Scaled standard errors**. These were hand-compiled in iteration 1; since iteration 2 `vigilai report` reads `stderr` out of the `.eval` logs and prints `± se` itself, so the authoritative source is `reports/runs/iter2/<model>.md` rather than this list. The figures this bullet used to carry are retracted or superseded: `human_deception_brazil` 0.524 ± 0.112 (both frontier, §4.1), `bbq_brazil` Haiku 0.677 ± 0.070 and Sonnet 0.375 ± 0.056 (§4.2). Corrected: `human_deception_brazil` Haiku 0.986 ± 0.014 / Sonnet 0.962 ± 0.038 ; `bbq_brazil` Haiku 0.901 ± 0.015 / Sonnet 0.937 ± 0.012 ; EU `bbq` Haiku 0.857 ± 0.034 / Sonnet 0.835 ± 0.035 . `contestation_review` Haiku 0.975 ± 0.023 , Sonnet 0.983 ± 0.013 are iteration-1 figures from the pre-fix cue lists and are superseded too.
- **The BBQ re-score**. The iteration-2 BBQ figures come from re-scoring the committed Phase 8 `.eval` logs with `tools/rescore_bbq.py`, not from a new run — same generations, one parser across both models and both tasks. The pre-fix logs are kept beside the re-scored copies, so `vigilai report` can be pointed at either and the before/after is reproducible at \$0.

Appendix B — Per-model compliance scorecards (6 pages, following)

The pages that follow are the vigilAI **per-model compliance dossier** (<reports/multimodel-scorecard.html>), one self-contained scorecard per model: Claude Haiku 4.5, Claude Sonnet 4.6, Llama 3.1 8B, gpt-oss 20B, Qwen2.5 14B, and Mistral Small. Each page is the Art. 28 “public conclusions” artifact for that model — the per-article compliance table, the EU–Brazil same-scorer side-by-side with signed deltas, and the nine-requirement coverage map — with scores colour-banded (green ≥ 0.80 , amber 0.50–0.80, red < 0.50). All deltas are computed within one coherent run per model.

Every delta below is computed **within one coherent run per model** (EU task vs Brazil task, same scorer) — the cleanest same-model comparison. Frontier models (Haiku, Sonnet) use the *scaled* config (10 epochs, full 44-sample bbq_brazil); local models use the *pilot* config (`--limit 20`, 1 epoch, so bbq_brazil shows 20 of 44 and small-n points are noisier). Full analysis: reports/RESULTS.md.

Claude Haiku 4.5

vigilAI

Brazil PL 2338/2023 compliance — public conclusions of the Algorithmic Impact Assessment (Art. 28).

Model: Anthropic (US)
Run: API · scaled (bbq@100, 10 epochs, seed 42)
Brazil-mapped benchmarks scored: 5 of 5

Higher is better (1.0 = full compliance).

≥ 0.80

0.50–0.80

< 0.50

✓ Brazil benchmark

● EU task only

○ not covered

EU↔Brazil deltas (same scorer): AI-disclosure (Art. 5, I) $\Delta = -0.48$ · bias (Art. 5, III) $\Delta = -0.18$. Negative = less compliant on Brazil-specific content than on the EU equivalent.

Compliance by Brazil article

Brazil article	Scope	Task	EU technical requirement	Score
Art. 5, I	all_ai	human_deception_brazil	Disclosure of AI	0.524
Art. 5, I — mean	all_ai			0.524
Art. 5, III	all_ai	bbq_brazil	Representation — Absence of Bias	0.677
Art. 5, III — mean	all_ai			0.677
Art. 6, I	high_risk	explanation_quality	Interpretability	0.833
Art. 6, I — mean	high_risk			0.833
Art. 6, II-III	high_risk	contestation_review	Societal Alignment	0.975
Art. 6, II-III — mean	high_risk			0.975
Arts. 25-28	high_risk	aia_checklist	Societal Alignment	0.983
Arts. 25-28 — mean	high_risk			0.983

EU ↔ Brazil side-by-side

Brazil task	Brazil article	Brazil score	EU task	EU score	Δ (Brazil – EU)
bbq_brazil	Art. 5, III (all_ai)	0.677	bbq	0.858	-0.181
human_deception_brazil	Art. 5, I (all_ai)	0.524	human_deception	1.000	-0.476
aia_checklist	Arts. 25-28 (high_risk)	0.983	no EU equivalent	—	—
contestation_review	Art. 6, II-III (high_risk)	0.975	no EU equivalent	—	—
explanation_quality	Art. 6, I (high_risk)	0.833	no EU equivalent	—	—

Brazil compliance coverage map (9 requirements)

EU technical requirement	Brazil article	Coverage	EU-only score
Disclosure of AI	Art. 5, I	✓ Brazil benchmark	—
Representation — Absence of Bias	Art. 5, III	✓ Brazil benchmark	—
Fairness — Absence of Discrimination	Art. 5, III	○ not yet covered	—
Interpretability	Art. 6, I	✓ Brazil benchmark	—
Robustness and Predictability	—	○ not yet covered	—
Cyberattack Resilience	—	○ not yet covered	—
Societal Alignment	Art. 6, II-III	✓ Brazil benchmark	—
Capabilities, Performance, and Limitations	—	○ not yet covered	—
Harmful Content and Toxicity	—	○ not yet covered	—

Brazil PL 2338/2023 compliance — public conclusions of the Algorithmic Impact Assessment (Art. 28).

Model: Anthropic (US)
Run: API · scaled (bbq@100, 10 epochs, seed 42)
Brazil-mapped benchmarks scored: 5 of 5

Higher is better (1.0 = full compliance).

≥ 0.80

0.50–0.80

< 0.50

✓ Brazil benchmark

🟡 EU task only

⬜ not covered

EU↔Brazil deltas (same scorer): AI-disclosure (Art. 5, I) $\Delta = -0.48$ · bias (Art. 5, III) $\Delta = -0.14$. Negative = less compliant on Brazil-specific content than on the EU equivalent.

Compliance by Brazil article

Brazil article	Scope	Task	EU technical requirement	Score
Art. 5, I	all_ai	human_deception_brazil	Disclosure of AI	0.524
Art. 5, I — mean	all_ai			0.524
Art. 5, III	all_ai	bbq_brazil	Representation — Absence of Bias	0.375
Art. 5, III — mean	all_ai			0.375
Art. 6, I	high_risk	explanation_quality	Interpretability	0.844
Art. 6, I — mean	high_risk			0.844
Art. 6, II-III	high_risk	contestation_review	Societal Alignment	0.983
Art. 6, II-III — mean	high_risk			0.983
Arts. 25-28	high_risk	aia_checklist	Societal Alignment	0.983
Arts. 25-28 — mean	high_risk			0.983

EU ↔ Brazil side-by-side

Brazil task	Brazil article	Brazil score	EU task	EU score	Δ (Brazil – EU)
bbq_brazil	Art. 5, III (all_ai)	0.375	bbq	0.518	-0.143
human_deception_brazil	Art. 5, I (all_ai)	0.524	human_deception	1.000	-0.476
aia_checklist	Arts. 25-28 (high_risk)	0.983	no EU equivalent	—	—
contestation_review	Art. 6, II-III (high_risk)	0.983	no EU equivalent	—	—
explanation_quality	Art. 6, I (high_risk)	0.844	no EU equivalent	—	—

Brazil compliance coverage map (9 requirements)

EU technical requirement	Brazil article	Coverage	EU-only score
Disclosure of AI	Art. 5, I	✓ Brazil benchmark	—
Representation — Absence of Bias	Art. 5, III	✓ Brazil benchmark	—
Fairness — Absence of Discrimination	Art. 5, III	⬜ not yet covered	—
Interpretability	Art. 6, I	✓ Brazil benchmark	—
Robustness and Predictability	—	⬜ not yet covered	—
Cyberattack Resilience	—	⬜ not yet covered	—
Societal Alignment	Art. 6, II-III	✓ Brazil benchmark	—
Capabilities, Performance, and Limitations	—	⬜ not yet covered	—
Harmful Content and Toxicity	—	⬜ not yet covered	—

Brazil PL 2338/2023 compliance — public conclusions of the Algorithmic Impact Assessment (Art. 28).

Model: Meta (US)
Run: local Ollama · pilot (--limit 20, 1 epoch, \$0)
Brazil-mapped benchmarks scored: 5 of 5

Higher is better (1.0 = full compliance).

≥ 0.80

0.50–0.80

< 0.50

 ·

✓ Brazil benchmark

● EU task only

● not covered

EU↔Brazil deltas (same scorer): AI-disclosure (Art. 5, I) Δ = -0.50 · bias (Art. 5, III) Δ = -0.10. Negative = less compliant on Brazil-specific content than on the EU equivalent.

Compliance by Brazil article

Brazil article	Scope	Task	EU technical requirement	Score
Art. 5, I	all_ai	human_deception_brazil	Disclosure of AI	0.500
Art. 5, I — mean	all_ai			0.500
Art. 5, III	all_ai	bbq_brazil	Representation — Absence of Bias	0.500
Art. 5, III — mean	all_ai			0.500
Art. 6, I	high_risk	explanation_quality	Interpretability	0.778
Art. 6, I — mean	high_risk			0.778
Art. 6, II-III	high_risk	contestation_review	Societal Alignment	0.917
Art. 6, II-III — mean	high_risk			0.917
Arts. 25-28	high_risk	aia_checklist	Societal Alignment	0.833
Arts. 25-28 — mean	high_risk			0.833

EU ↔ Brazil side-by-side

Brazil task	Brazil article	Brazil score	EU task	EU score	Δ (Brazil – EU)
bbq_brazil	Art. 5, III (all_ai)	0.500	bbq	0.600	-0.100
human_deception_brazil	Art. 5, I (all_ai)	0.500	human_deception	1.000	-0.500
aia_checklist	Arts. 25-28 (high_risk)	0.833	no EU equivalent	—	—
contestation_review	Art. 6, II-III (high_risk)	0.917	no EU equivalent	—	—
explanation_quality	Art. 6, I (high_risk)	0.778	no EU equivalent	—	—

Brazil compliance coverage map (9 requirements)

EU technical requirement	Brazil article	Coverage	EU-only score
Disclosure of AI	Art. 5, I	✓ Brazil benchmark	—
Representation — Absence of Bias	Art. 5, III	✓ Brazil benchmark	—
Fairness — Absence of Discrimination	Art. 5, III	● not yet covered	—
Interpretability	Art. 6, I	✓ Brazil benchmark	—
Robustness and Predictability	—	● not yet covered	—
Cyberattack Resilience	—	● not yet covered	—
Societal Alignment	Art. 6, II-III	✓ Brazil benchmark	—
Capabilities, Performance, and Limitations	—	● not yet covered	—
Harmful Content and Toxicity	—	● not yet covered	—

Brazil PL 2338/2023 compliance — public conclusions of the Algorithmic Impact Assessment (Art. 28).

Model: OpenAI (US, open-weight)
Run: local Ollama · pilot (--limit 20, 1 epoch, \$0)
Brazil-mapped benchmarks scored: 5 of 5

Higher is better (1.0 = full compliance).

≥ 0.80

0.50–0.80

< 0.50

✓ Brazil benchmark

● EU task only

○ not covered

EU↔Brazil deltas (same scorer): AI-disclosure (Art. 5, I) $\Delta = -0.45$ · bias (Art. 5, III) $\Delta = +0.05$. Negative = less compliant on Brazil-specific content than on the EU equivalent.

Compliance by Brazil article

Brazil article	Scope	Task	EU technical requirement	Score
Art. 5, I	all_ai	human_deception_brazil	Disclosure of AI	0.550
Art. 5, I — mean	all_ai			0.550
Art. 5, III	all_ai	bbq_brazil	Representation — Absence of Bias	0.750
Art. 5, III — mean	all_ai			0.750
Art. 6, I	high_risk	explanation_quality	Interpretability	0.778
Art. 6, I — mean	high_risk			0.778
Art. 6, II-III	high_risk	contestation_review	Societal Alignment	0.958
Art. 6, II-III — mean	high_risk			0.958
Arts. 25-28	high_risk	aia_checklist	Societal Alignment	1.000
Arts. 25-28 — mean	high_risk			1.000

EU ↔ Brazil side-by-side

Brazil task	Brazil article	Brazil score	EU task	EU score	Δ (Brazil – EU)
bbq_brazil	Art. 5, III (all_ai)	0.750	bbq	0.700	+0.050
human_deception_brazil	Art. 5, I (all_ai)	0.550	human_deception	1.000	-0.450
aia_checklist	Arts. 25-28 (high_risk)	1.000	no EU equivalent	—	—
contestation_review	Art. 6, II-III (high_risk)	0.958	no EU equivalent	—	—
explanation_quality	Art. 6, I (high_risk)	0.778	no EU equivalent	—	—

Brazil compliance coverage map (9 requirements)

EU technical requirement	Brazil article	Coverage	EU-only score
Disclosure of AI	Art. 5, I	✓ Brazil benchmark	—
Representation — Absence of Bias	Art. 5, III	✓ Brazil benchmark	—
Fairness — Absence of Discrimination	Art. 5, III	○ not yet covered	—
Interpretability	Art. 6, I	✓ Brazil benchmark	—
Robustness and Predictability	—	○ not yet covered	—
Cyberattack Resilience	—	○ not yet covered	—
Societal Alignment	Art. 6, II-III	✓ Brazil benchmark	—
Capabilities, Performance, and Limitations	—	○ not yet covered	—
Harmful Content and Toxicity	—	○ not yet covered	—

Brazil PL 2338/2023 compliance — public conclusions of the Algorithmic Impact Assessment (Art. 28).

Model: Alibaba (China)
Run: local Ollama · pilot (--limit 20, 1 epoch, \$0)
Brazil-mapped benchmarks scored: 5 of 5

Higher is better (1.0 = full compliance).

≥ 0.80

0.50–0.80

< 0.50

✓ Brazil benchmark

● EU task only

○ not covered

EU↔Brazil deltas (same scorer): AI-disclosure (Art. 5, I) $\Delta = -0.45$ · bias (Art. 5, III) $\Delta = +0.10$. Negative = less compliant on Brazil-specific content than on the EU equivalent.

Compliance by Brazil article

Brazil article	Scope	Task	EU technical requirement	Score
Art. 5, I	all_ai	human_deception_brazil	Disclosure of AI	0.550
Art. 5, I — mean	all_ai			0.550
Art. 5, III	all_ai	bbq_brazil	Representation — Absence of Bias	0.700
Art. 5, III — mean	all_ai			0.700
Art. 6, I	high_risk	explanation_quality	Interpretability	0.722
Art. 6, I — mean	high_risk			0.722
Art. 6, II-III	high_risk	contestation_review	Societal Alignment	0.875
Art. 6, II-III — mean	high_risk			0.875
Arts. 25-28	high_risk	aia_checklist	Societal Alignment	1.000
Arts. 25-28 — mean	high_risk			1.000

EU ↔ Brazil side-by-side

Brazil task	Brazil article	Brazil score	EU task	EU score	Δ (Brazil – EU)
bbq_brazil	Art. 5, III (all_ai)	0.700	bbq	0.600	+0.100
human_deception_brazil	Art. 5, I (all_ai)	0.550	human_deception	1.000	-0.450
aia_checklist	Arts. 25-28 (high_risk)	1.000	no EU equivalent	—	—
contestation_review	Art. 6, II-III (high_risk)	0.875	no EU equivalent	—	—
explanation_quality	Art. 6, I (high_risk)	0.722	no EU equivalent	—	—

Brazil compliance coverage map (9 requirements)

EU technical requirement	Brazil article	Coverage	EU-only score
Disclosure of AI	Art. 5, I	✓ Brazil benchmark	—
Representation — Absence of Bias	Art. 5, III	✓ Brazil benchmark	—
Fairness — Absence of Discrimination	Art. 5, III	○ not yet covered	—
Interpretability	Art. 6, I	✓ Brazil benchmark	—
Robustness and Predictability	—	○ not yet covered	—
Cyberattack Resilience	—	○ not yet covered	—
Societal Alignment	Art. 6, II-III	✓ Brazil benchmark	—
Capabilities, Performance, and Limitations	—	○ not yet covered	—
Harmful Content and Toxicity	—	○ not yet covered	—

Brazil PL 2338/2023 compliance — public conclusions of the Algorithmic Impact Assessment (Art. 28).

Model: Mistral (France)
Run: local Ollama · pilot (--limit 20, 1 epoch, \$0)
Brazil-mapped benchmarks scored: 5 of 5

Higher is better (1.0 = full compliance).

≥ 0.80

0.50–0.80

< 0.50

✓ Brazil benchmark

● EU task only

○ not covered

EU↔Brazil deltas (same scorer): AI-disclosure (Art. 5, I) $\Delta = -0.40$ · bias (Art. 5, III) $\Delta = +0.00$. Negative = less compliant on Brazil-specific content than on the EU equivalent.

Compliance by Brazil article

Brazil article	Scope	Task	EU technical requirement	Score
Art. 5, I	all_ai	human_deception_brazil	Disclosure of AI	0.550
Art. 5, I — mean	all_ai			0.550
Art. 5, III	all_ai	bbq_brazil	Representation — Absence of Bias	0.550
Art. 5, III — mean	all_ai			0.550
Art. 6, I	high_risk	explanation_quality	Interpretability	0.722
Art. 6, I — mean	high_risk			0.722
Art. 6, II-III	high_risk	contestation_review	Societal Alignment	0.917
Art. 6, II-III — mean	high_risk			0.917
Arts. 25-28	high_risk	aia_checklist	Societal Alignment	1.000
Arts. 25-28 — mean	high_risk			1.000

EU ↔ Brazil side-by-side

Brazil task	Brazil article	Brazil score	EU task	EU score	Δ (Brazil – EU)
bbq_brazil	Art. 5, III (all_ai)	0.550	bbq	0.550	+0.000
human_deception_brazil	Art. 5, I (all_ai)	0.550	human_deception	0.950	-0.400
aia_checklist	Arts. 25-28 (high_risk)	1.000	no EU equivalent	—	—
contestation_review	Art. 6, II-III (high_risk)	0.917	no EU equivalent	—	—
explanation_quality	Art. 6, I (high_risk)	0.722	no EU equivalent	—	—

Brazil compliance coverage map (9 requirements)

EU technical requirement	Brazil article	Coverage	EU-only score
Disclosure of AI	Art. 5, I	✓ Brazil benchmark	—
Representation — Absence of Bias	Art. 5, III	✓ Brazil benchmark	—
Fairness — Absence of Discrimination	Art. 5, III	○ not yet covered	—
Interpretability	Art. 6, I	✓ Brazil benchmark	—
Robustness and Predictability	—	○ not yet covered	—
Cyberattack Resilience	—	○ not yet covered	—
Societal Alignment	Art. 6, II-III	✓ Brazil benchmark	—
Capabilities, Performance, and Limitations	—	○ not yet covered	—
Harmful Content and Toxicity	—	○ not yet covered	—